

# Deep Learning in MIR: **Demystifying the Dark Art**

ISMIR Tutorial Session

Curitiba, Brazil  
4 November 2013



Erik M. Schmidt  
eschmidt@pandora.com



Philipe Hamel  
hamelphi@google.com



Eric J. Humphrey  
ejh333@nyu.edu

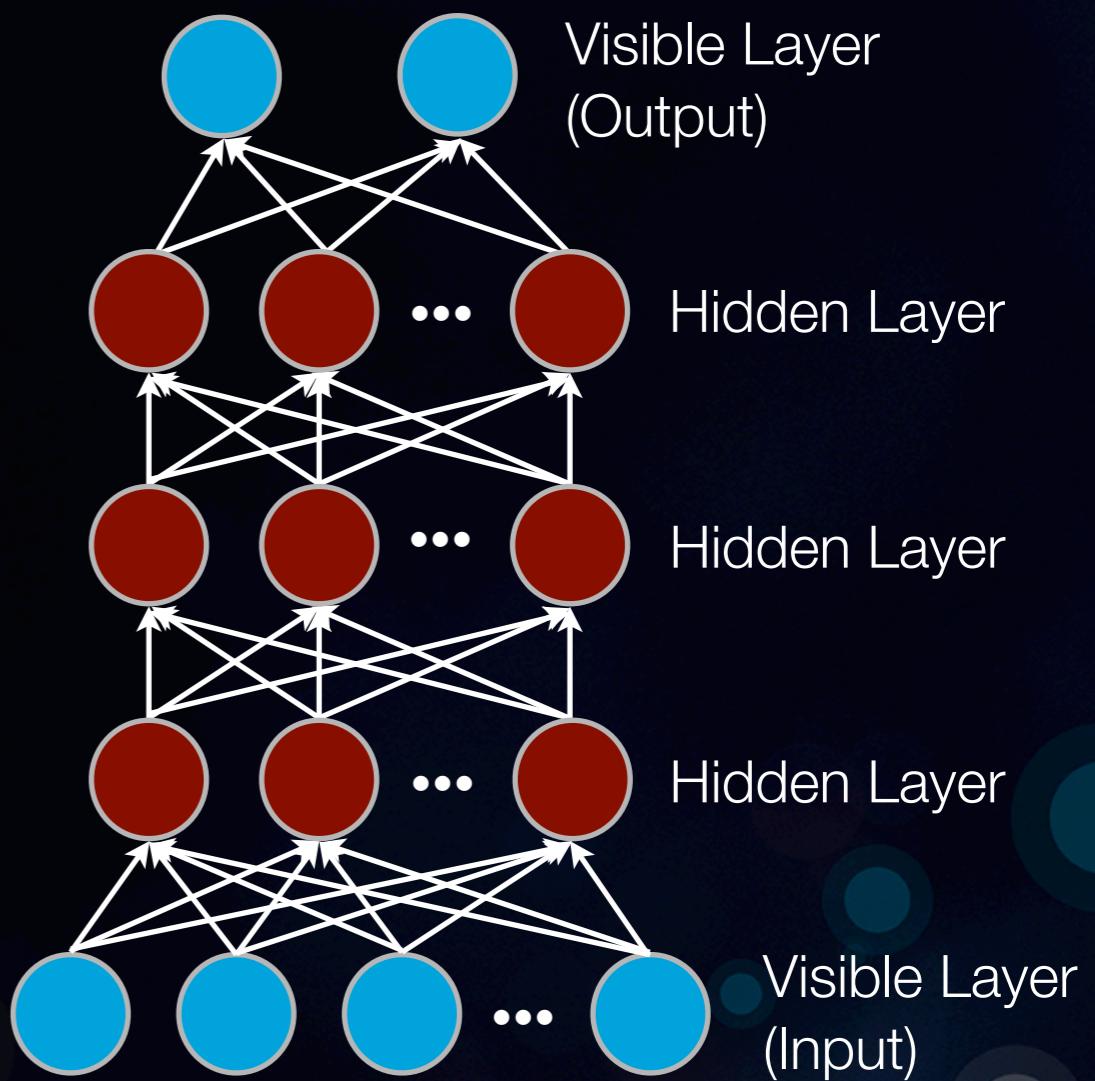
# Deep Learning

What exactly does that mean?

- A deep learning model simply is one which contains multiple layers
- A layer can be any simple transformation:

$$f(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b$$

- Generally contain non-linearities
- **Many names:** Artificial neural nets, multi-layer perceptrons, convolutional neural networks, deep belief networks, etc



# Deep Learning in MIR

## Tutorial Overview

Erik Schmidt  
Pandora Media, Inc

- History of deep learning in AI
  - Perceptrons, neural nets, SVMs, convolutional neural nets, restricted Boltzman machines, greedy-wise pre-training, dropout
- Deep learning in music emotion recognition
  - Learning emotion-based representations from raw magnitude spectra
  - Learning melody and rhythm features with deep belief networks

# Deep Learning in MIR

## Tutorial Overview

Philippe Hamel  
Google, Inc

- The state-of-the-art in deep learning
  - Pre-training vs. initialization
  - Dropout
  - Activation functions
  - Structure (convolution, recurrence)
- Deep learning in MIR

# Deep Learning in MIR

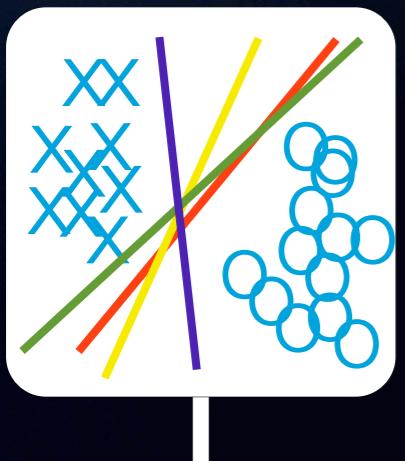
## Tutorial Overview

Eric Humphrey  
NYU MARL

- Deep learning through a signal processing lens
  - What “is” deep learning
  - How you are already doing it (kinda)
  - How you might do it intentionally
  - Some tips, tricks, insight, and advice
  - A few thoughts for the future

# Deep Learning

A long, long time ago...

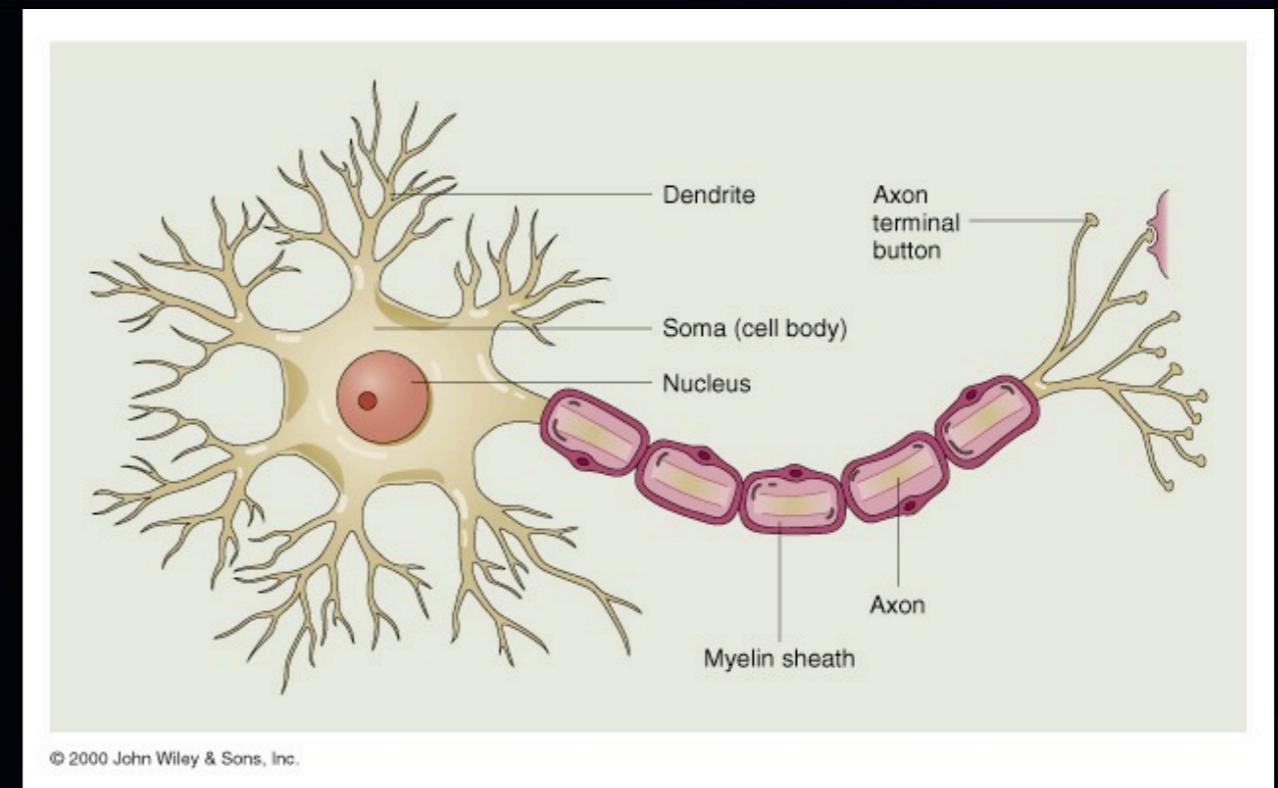


1959  
Perceptrons

# Perceptrons

## Neural inspired models

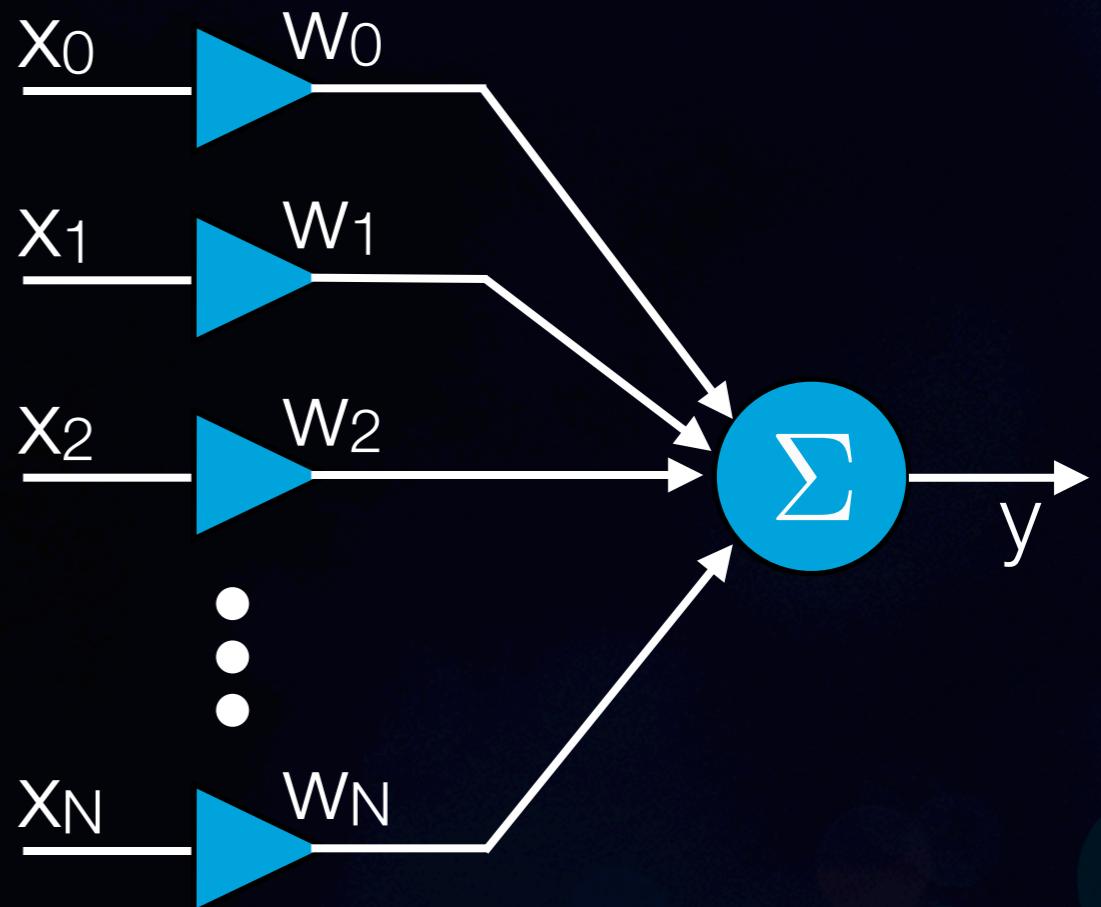
- Dendrite: input terminals
- Cell body: signals combined
- Axon: output cable
- Synapse: interface to other neurons



# Perceptrons

## Neural inspired models

- The perceptron is a model inspired by the neuron
- One of the earliest machine learning models
- Later we will see these as a building block for deep learning



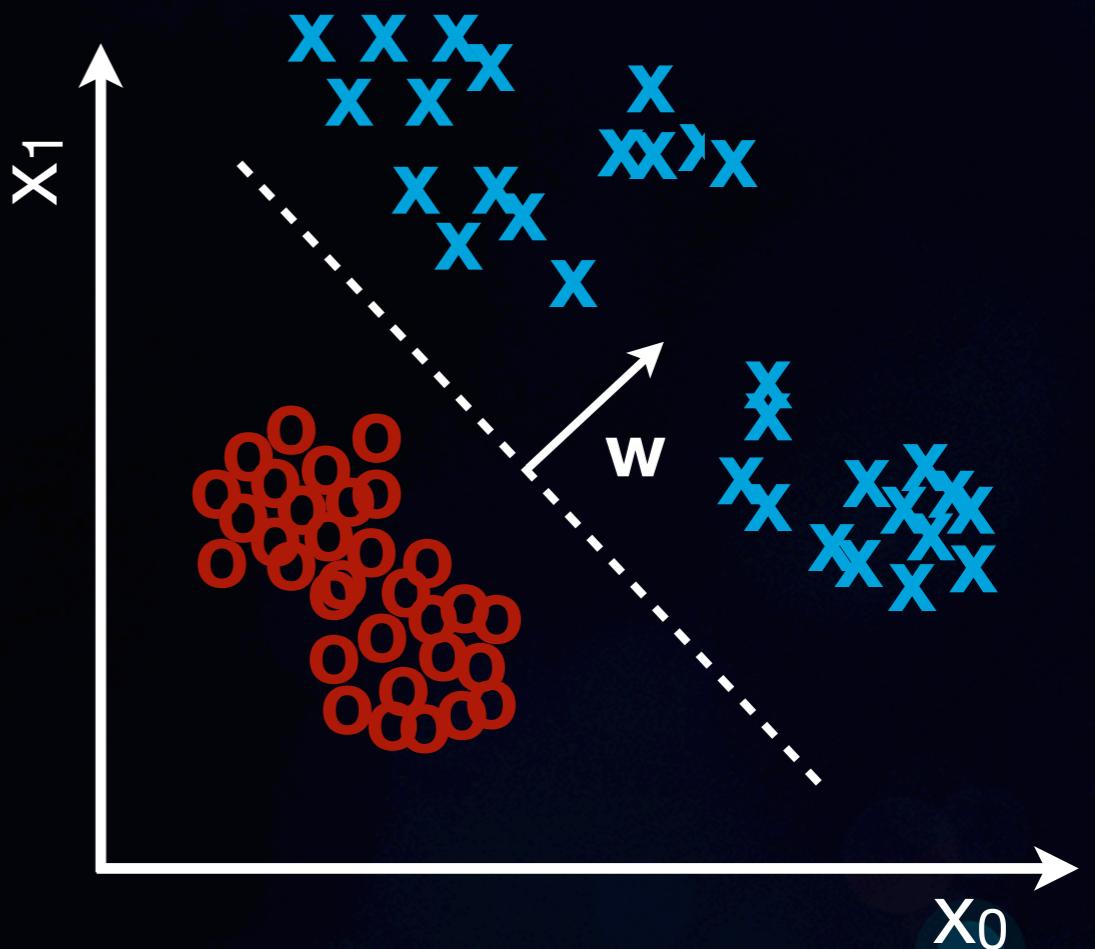
$$f(x) = x_0w_0 + \dots + x_Nw_N$$

# Perceptrons

## Neural inspired models

- Include intercept term  $b$
- Use inner product notation
- Decision function that's very easy to evaluate on a computer

$$f(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b$$



# Perceptron Algorithm

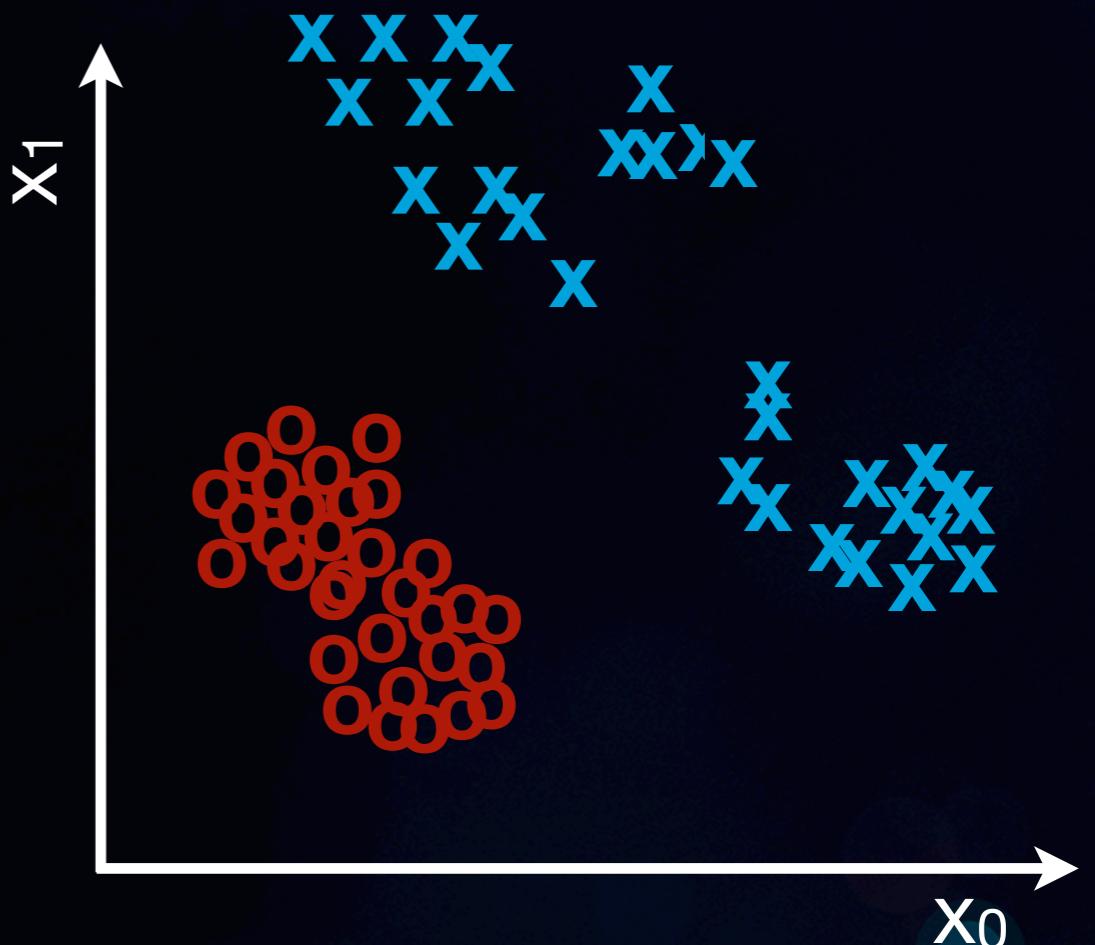
Fitting the model to data

- Input data:
  - Training data  $\mathbf{x}$
  - Labels  $y = +/-1$
- Learning Objective:

$$\text{if } y_i = 1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b > 0$$

$$\text{if } y_i = -1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b < 0$$

$$\text{for } \forall y \\ y_i(\langle \mathbf{w}, \mathbf{x}_i \rangle + b) > 0$$



# Perceptron Algorithm

Fitting the model to data

initialize  $w, b = 0$

repeat

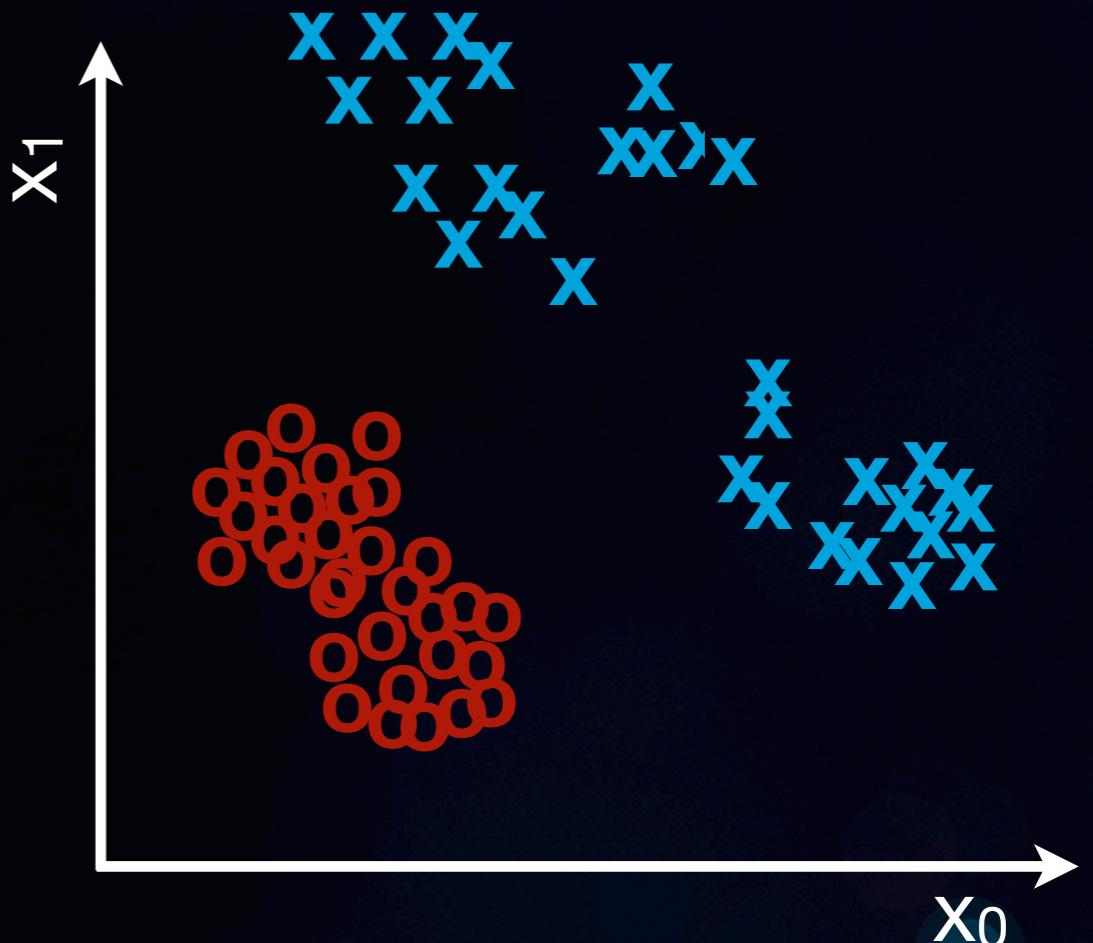
Pick  $(x_i, y_i)$  from data

if  $y_i(\langle w, x_i \rangle + b) \leq 0$  then

$w' = w + y_i x_i$

$b' = b + y_i$

until  $y_i(\langle w, x_i \rangle + b) > 0$  for  $\forall i$



# Perceptron Algorithm

Fitting the model to data

initialize  $w, b = 0$

repeat

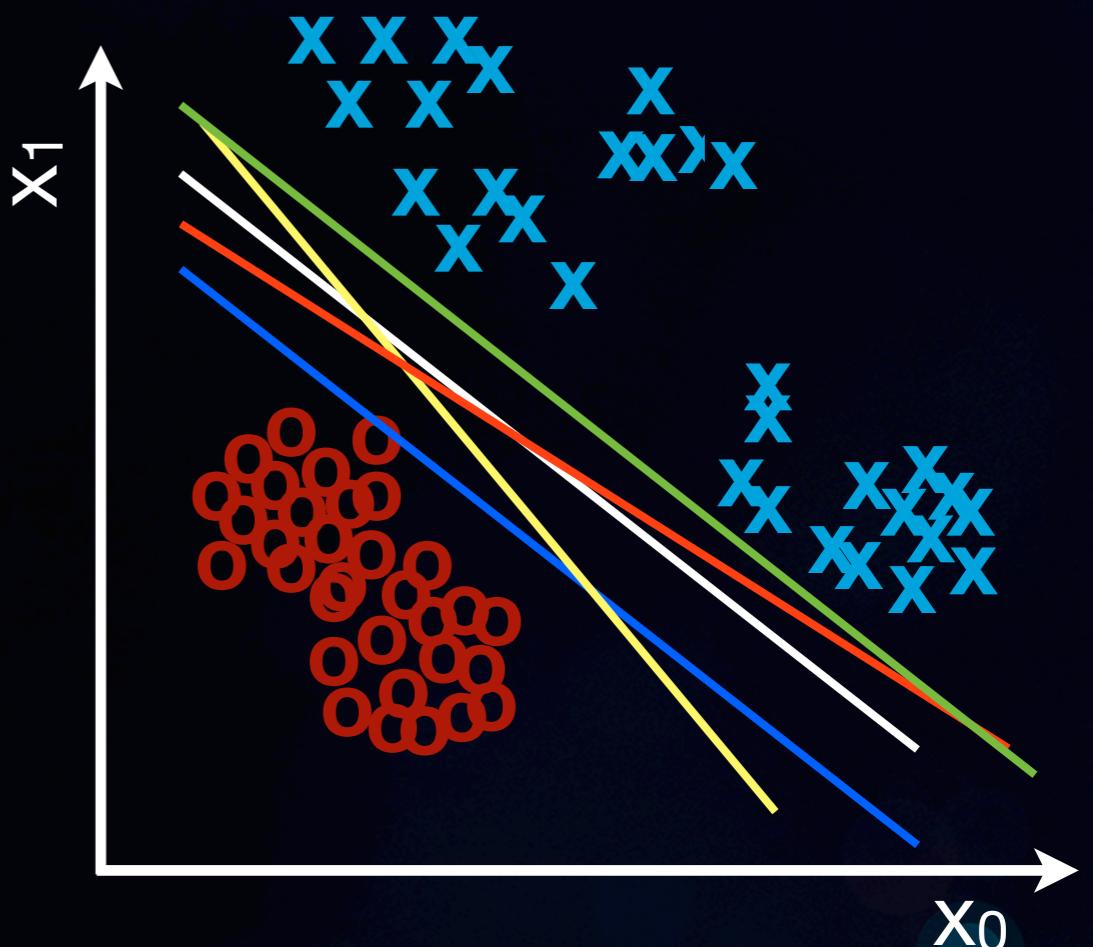
Pick  $(x_i, y_i)$  from data

if  $y_i(\langle w, x_i \rangle + b) \leq 0$  then

$w' = w + y_i x_i$

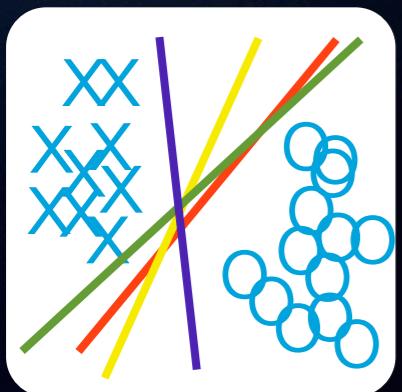
$b' = b + y_i$

until  $y_i(\langle w, x_i \rangle + b) > 0$  for  $\forall i$

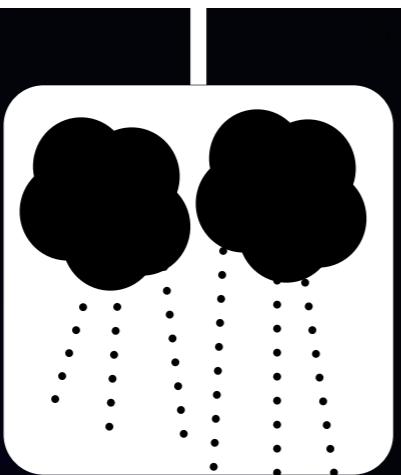


# Deep Learning

A long, long time ago...



1959  
Perceptrons

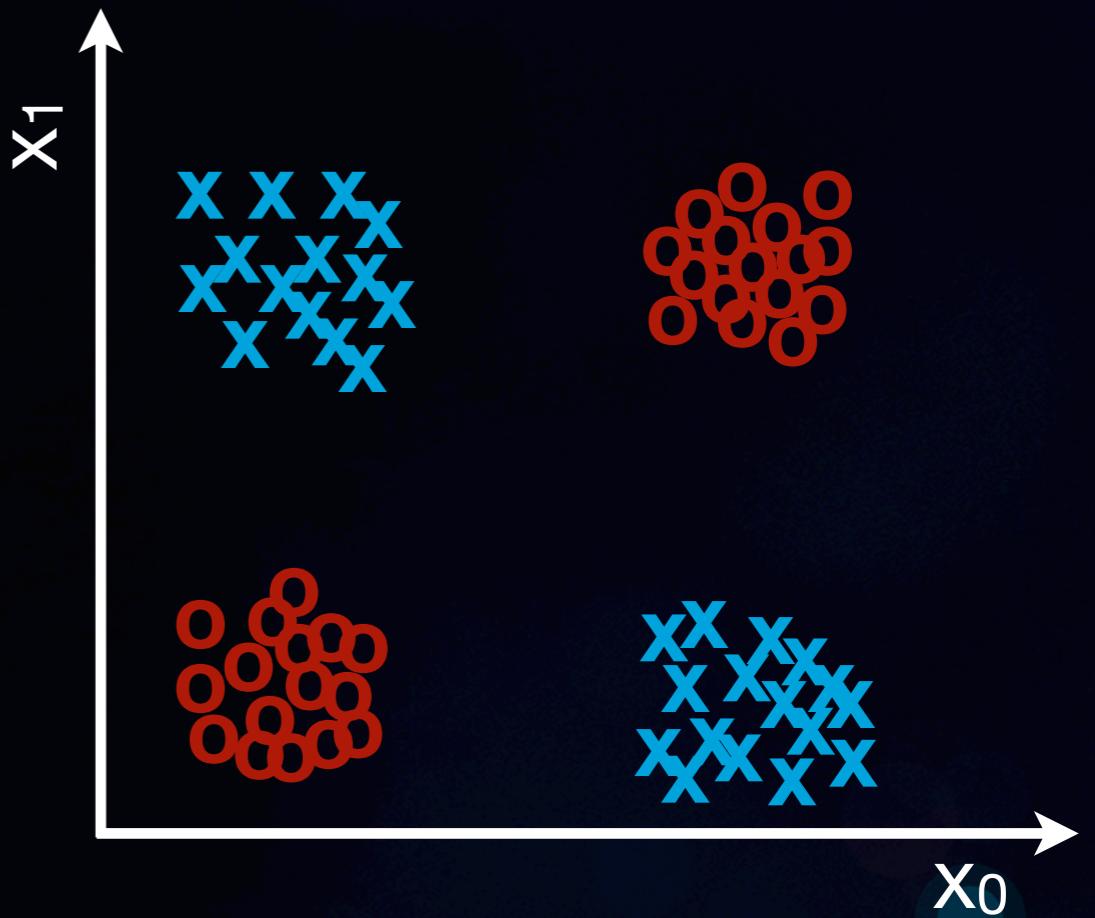


AI Winter  
1969

# XOR Problem

A major limitation

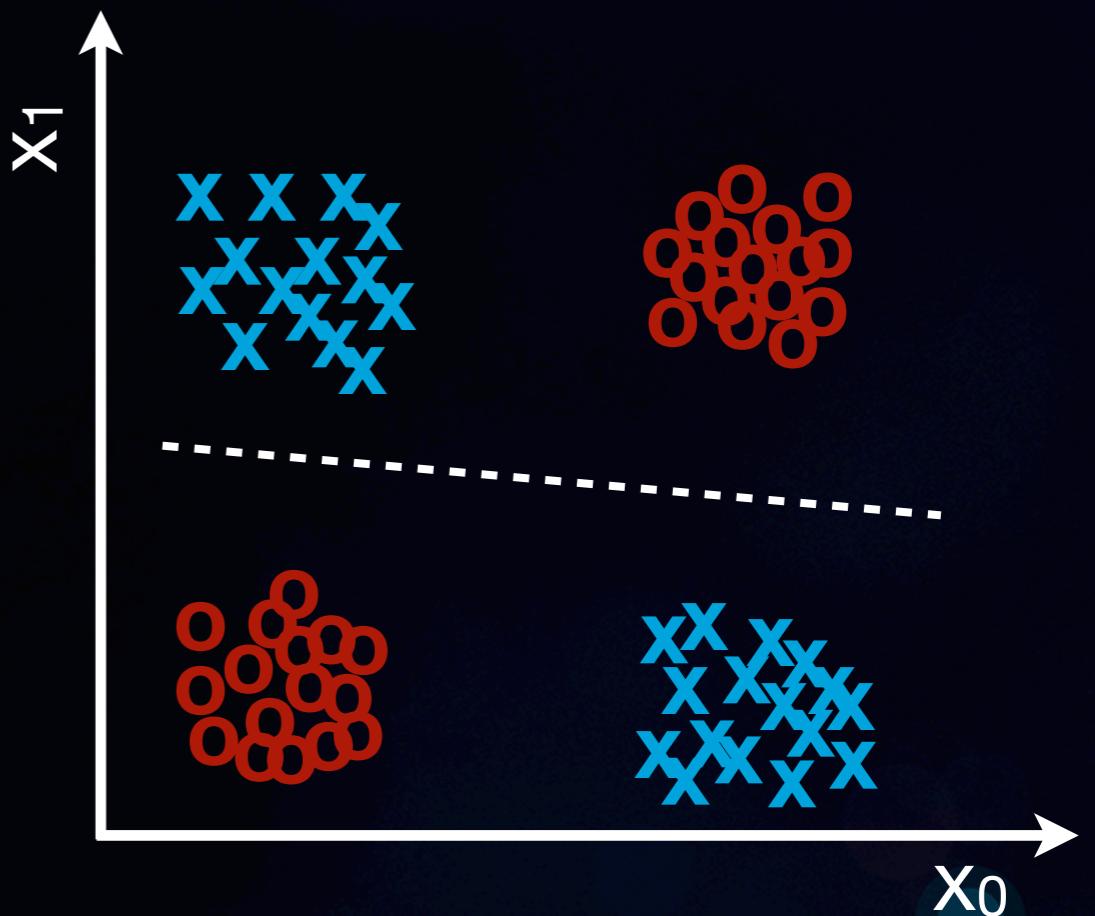
- The perceptron algorithm can only learn a linear discriminant
- The inability to learn an XOR function was demonstrated by Marvin Minsky



# XOR Problem

A major limitation

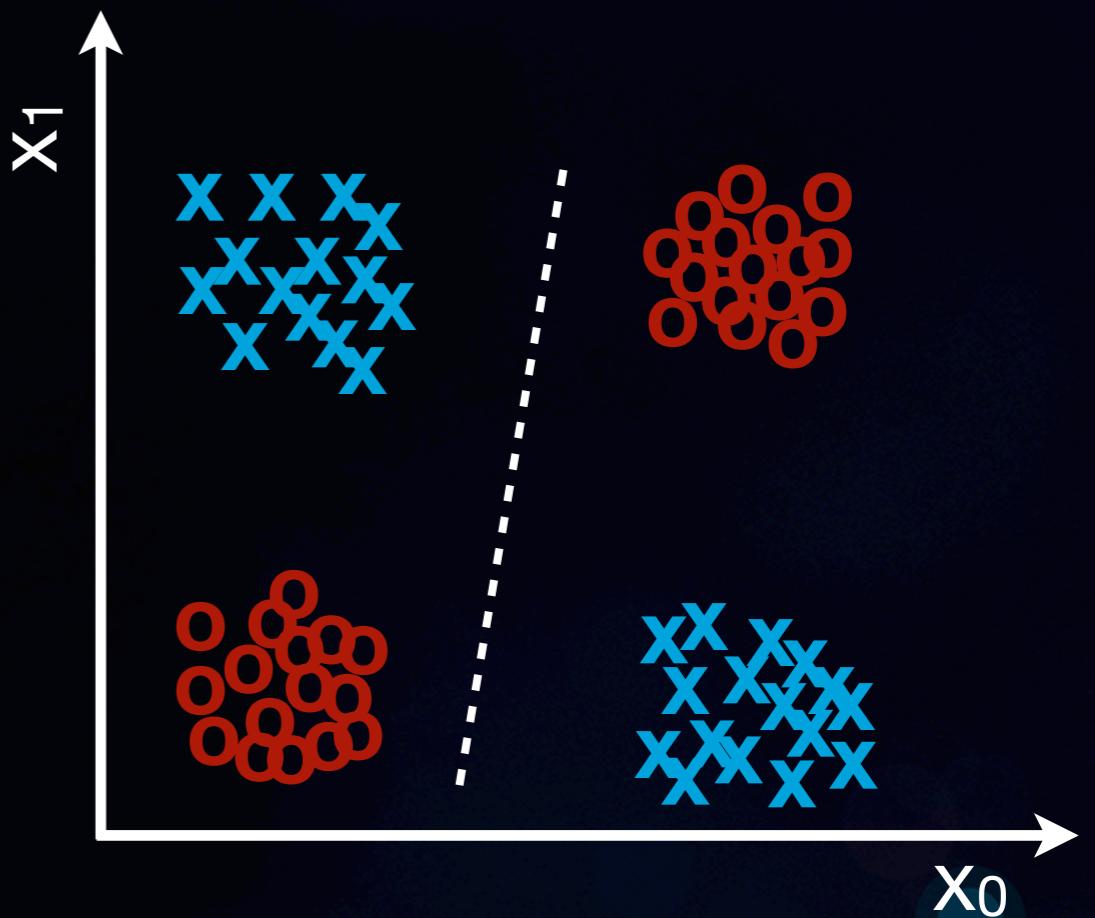
- The perceptron algorithm can only learn a linear discriminant
- The inability to learn an XOR function was demonstrated by Marvin Minsky



# XOR Problem

A major limitation

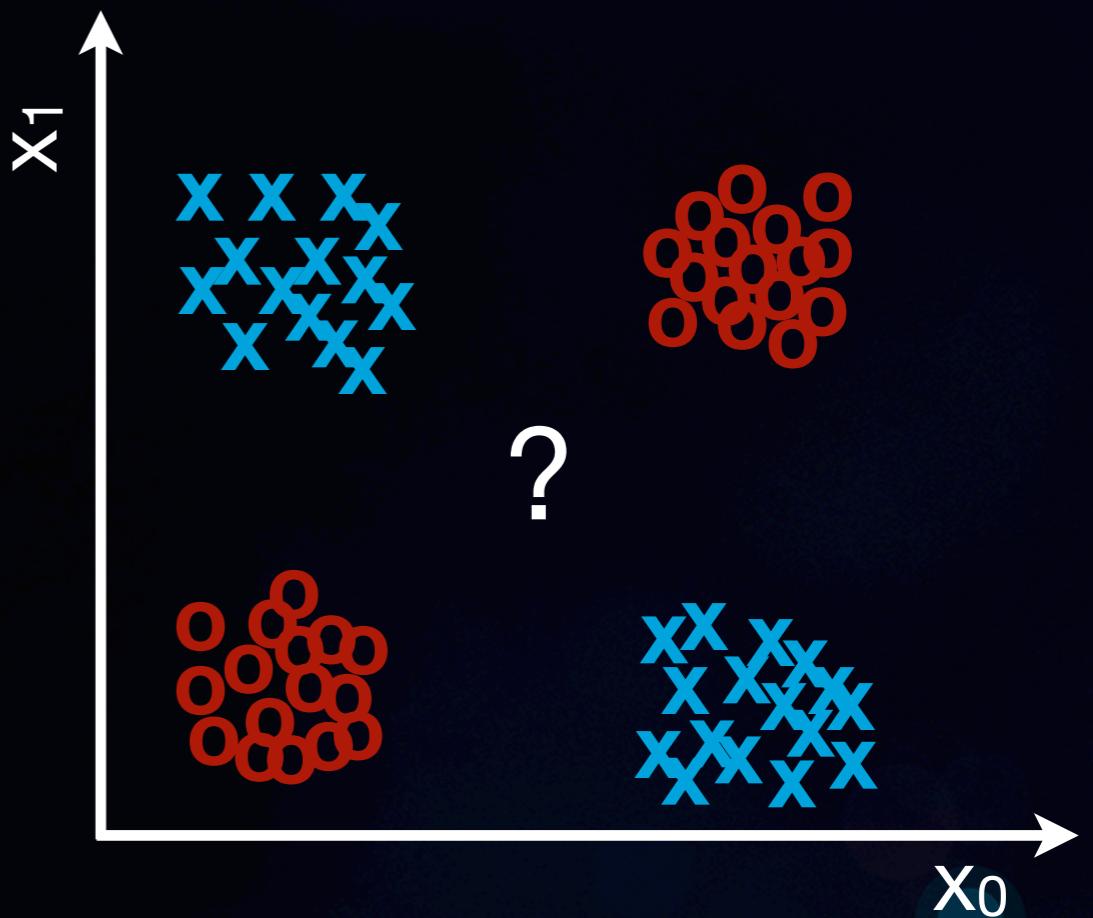
- The perceptron algorithm can only learn a linear discriminant
- The inability to learn an XOR function was demonstrated by Marvin Minsky



# XOR Problem

A major limitation

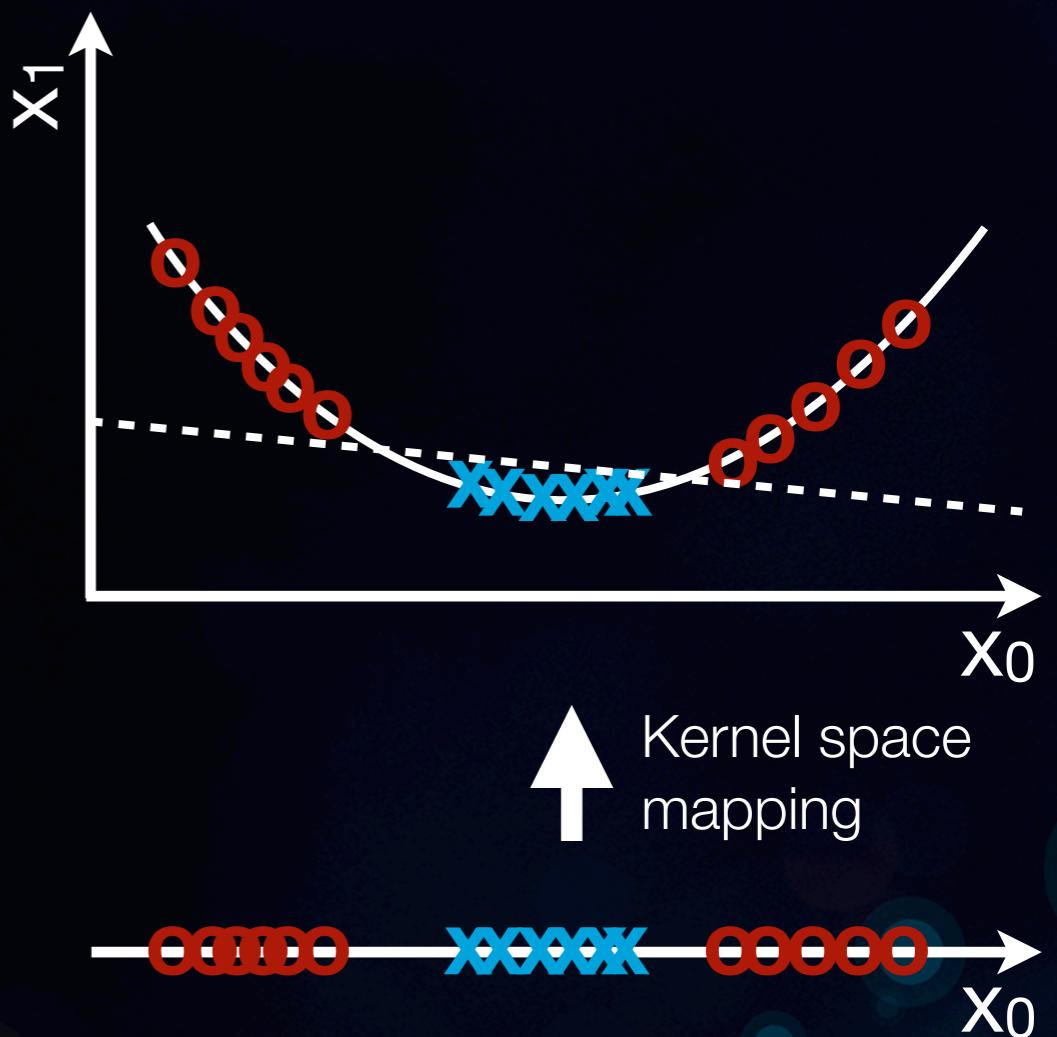
- The perceptron algorithm can only learn a linear discriminant
- The inability to learn an XOR function was demonstrated by Marvin Minsky



# Kernel Trick

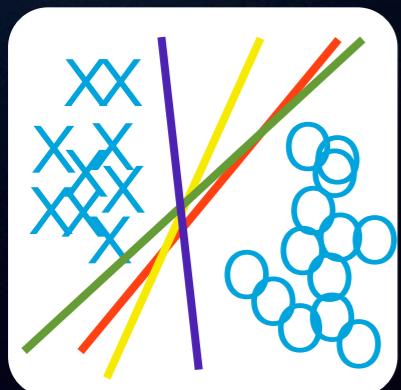
Solving non-linear problems

- Project data into higher dimensional space
- Employ non-linear transformation
- Obtain linear separability

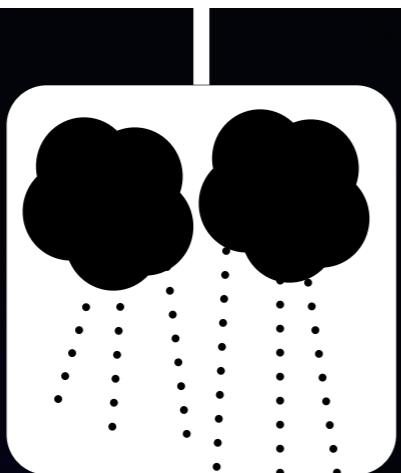


# Deep Learning

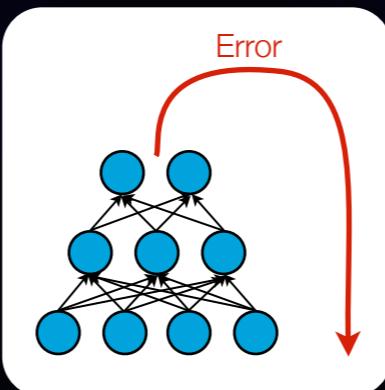
A long, long time ago...



1959  
Perceptrons



AI Winter  
1969

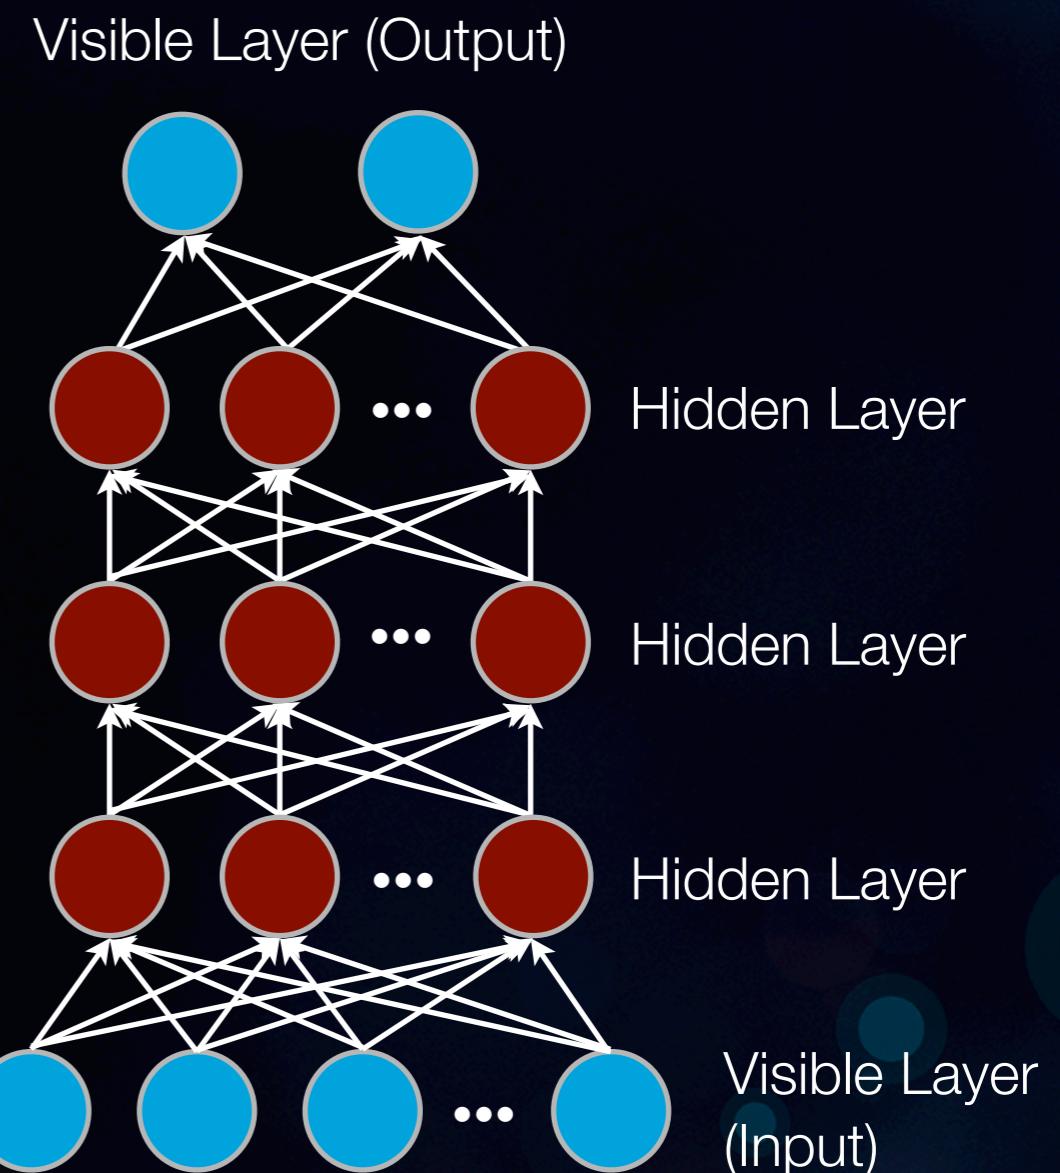


1986  
Artificial Neural  
Networks

# Artificial Neural Networks

## Adding in multiple layers

- In the 1980s researchers began making progress on deeper representations
- Build up larger models inspired by neural interaction in human brain

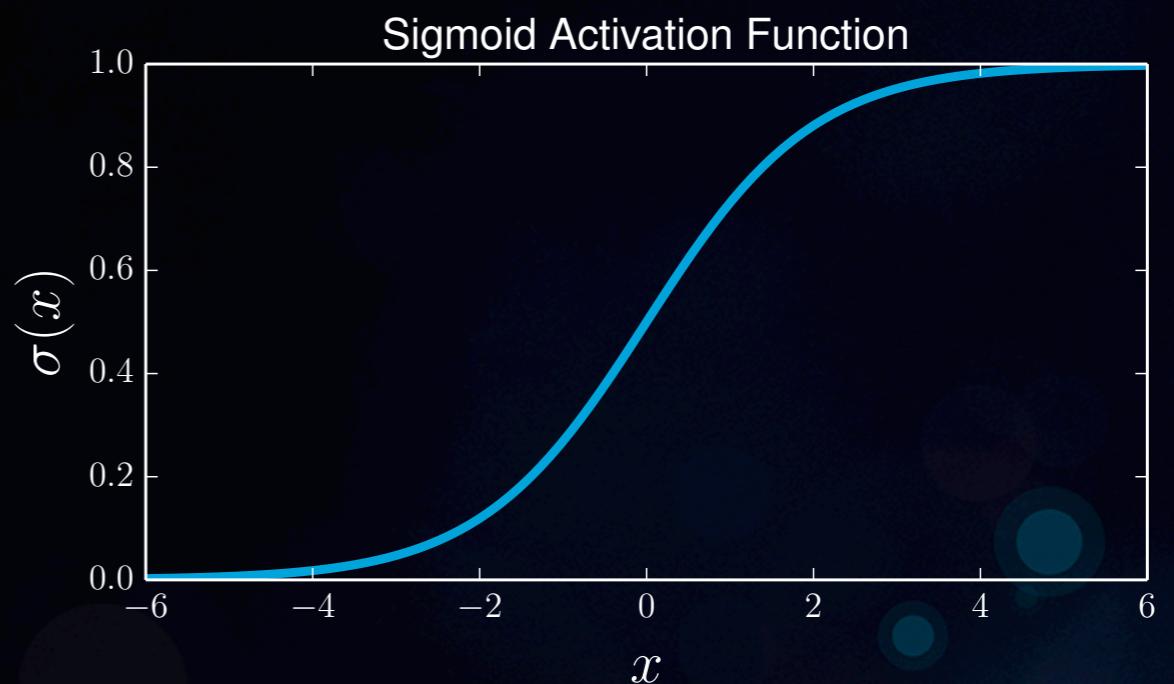


# Artificial Neural Networks

## Adding in multiple layers

- Each layer output applies a non-linearity or “activation”
- Most common is sigmoid activation
- Allows learning of non-linear relationships
- Without non-linearities deep network could be represented in a single layer

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



# How Do We Train?

These models are complicated!

1-layer Model:

$$f(\mathbf{x}) = \sigma (\langle \mathbf{w}_0, \mathbf{x} \rangle + b_0)$$

2 parameters to fit...

2-layer Model:

$$f(\mathbf{x}) = \sigma (\langle \mathbf{w}_1, \sigma (\langle \mathbf{w}_0, \mathbf{x} \rangle + b_0) \rangle + b_1)$$

4 parameters to fit...

3-layer Model:

$$f(\mathbf{x}) = \sigma (\langle \mathbf{w}_2, \sigma (\langle \mathbf{w}_1, \sigma (\langle \mathbf{w}_0, \mathbf{x} \rangle + b_0) \rangle + b_1) \rangle + b_2)$$

6 parameters to fit!!!

# Backpropagation

An algorithm for deep neural nets

Simple 1-layer Model:

$$\hat{y} = \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b)$$

Simple cost function:

$$\text{cost} = \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

Compute partials:

$$\partial \mathbf{w} = \frac{\partial}{\partial \mathbf{w}} \text{cost} = \frac{\partial}{\partial \mathbf{w}} \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

$$\partial b = \frac{\partial}{\partial b} \text{cost} = \frac{\partial}{\partial b} \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

initialize  $\theta$  (small random values)

repeat

Pick  $(x_i, y_i)$  from data

Compute model error

Compute:  $\Delta(\theta^{(t)})$

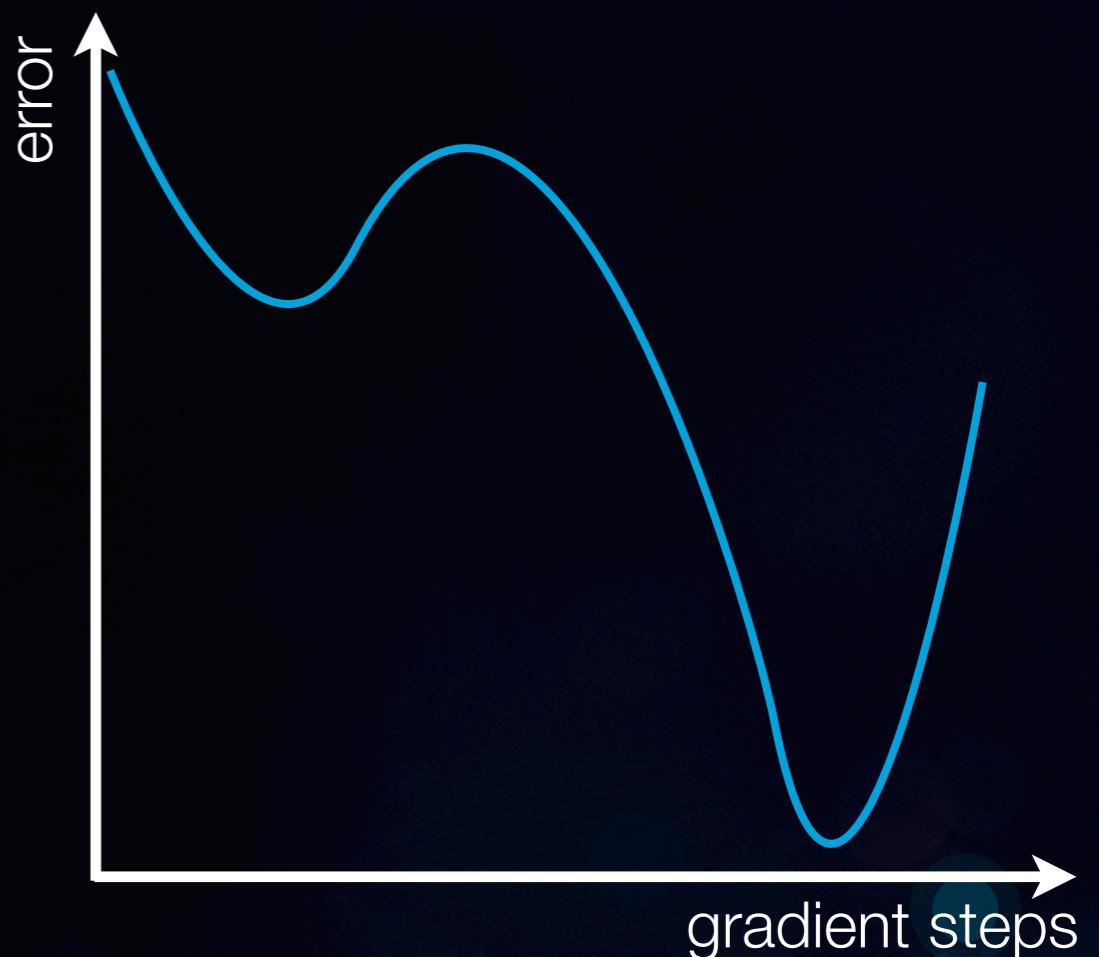
Update:  $\theta^{(t+1)} = \theta^{(t)} - \epsilon \Delta E(\theta^{(t)})$

until stopping criteria satisfied

# Backpropagation

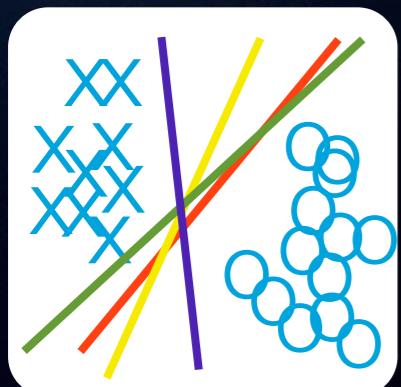
## Issues

- Optimization guarantees a local minima, not global
- Small datasets produce noisy gradients
  - Lots of local minima
- Limited computational resources

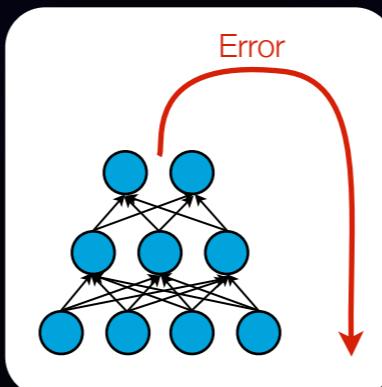


# Deep Learning

A long, long time ago...

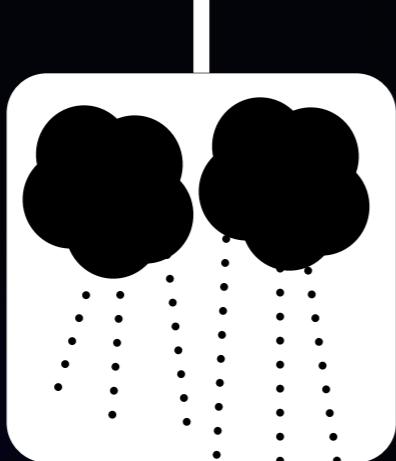


AI Winter  
1969

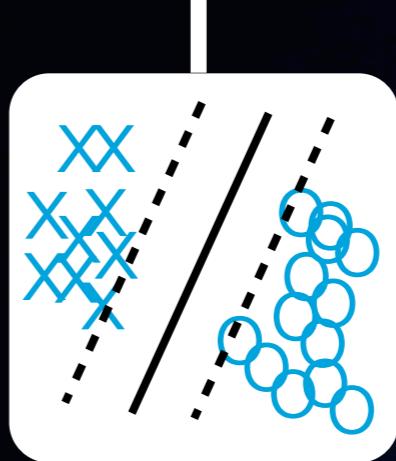


Support Vector  
Machines  
1995

1959  
Perceptrons



1986  
Artificial Neural  
Networks



# Support Vector Machines

A free lunch?

New constraints:

$$\text{if } y_i = 1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b \geq 1$$

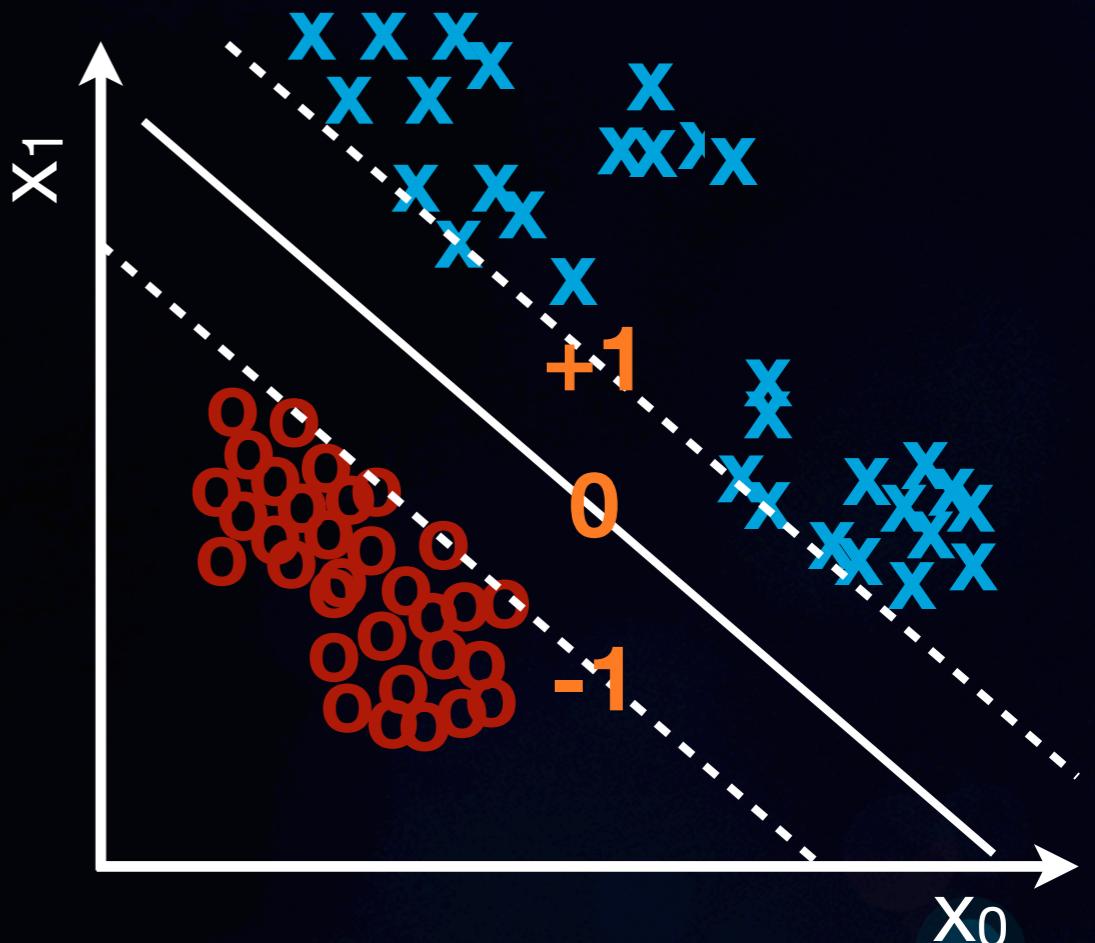
$$\text{if } y_i = -1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b \leq -1$$

$$\text{for } \forall y \\ y_i(\langle \mathbf{w}, \mathbf{x}_i \rangle + b) \geq 1$$

Optimization Function:

$$\text{minimize } \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{subject to } y_i(\langle \mathbf{w}, \mathbf{x}_i \rangle + b) - 1 \geq 0 \text{ for all } 1 \leq i \leq m$$



# Support Vector Machines

A free lunch?

New constraints:

$$\text{if } y_i = 1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b \geq 1$$

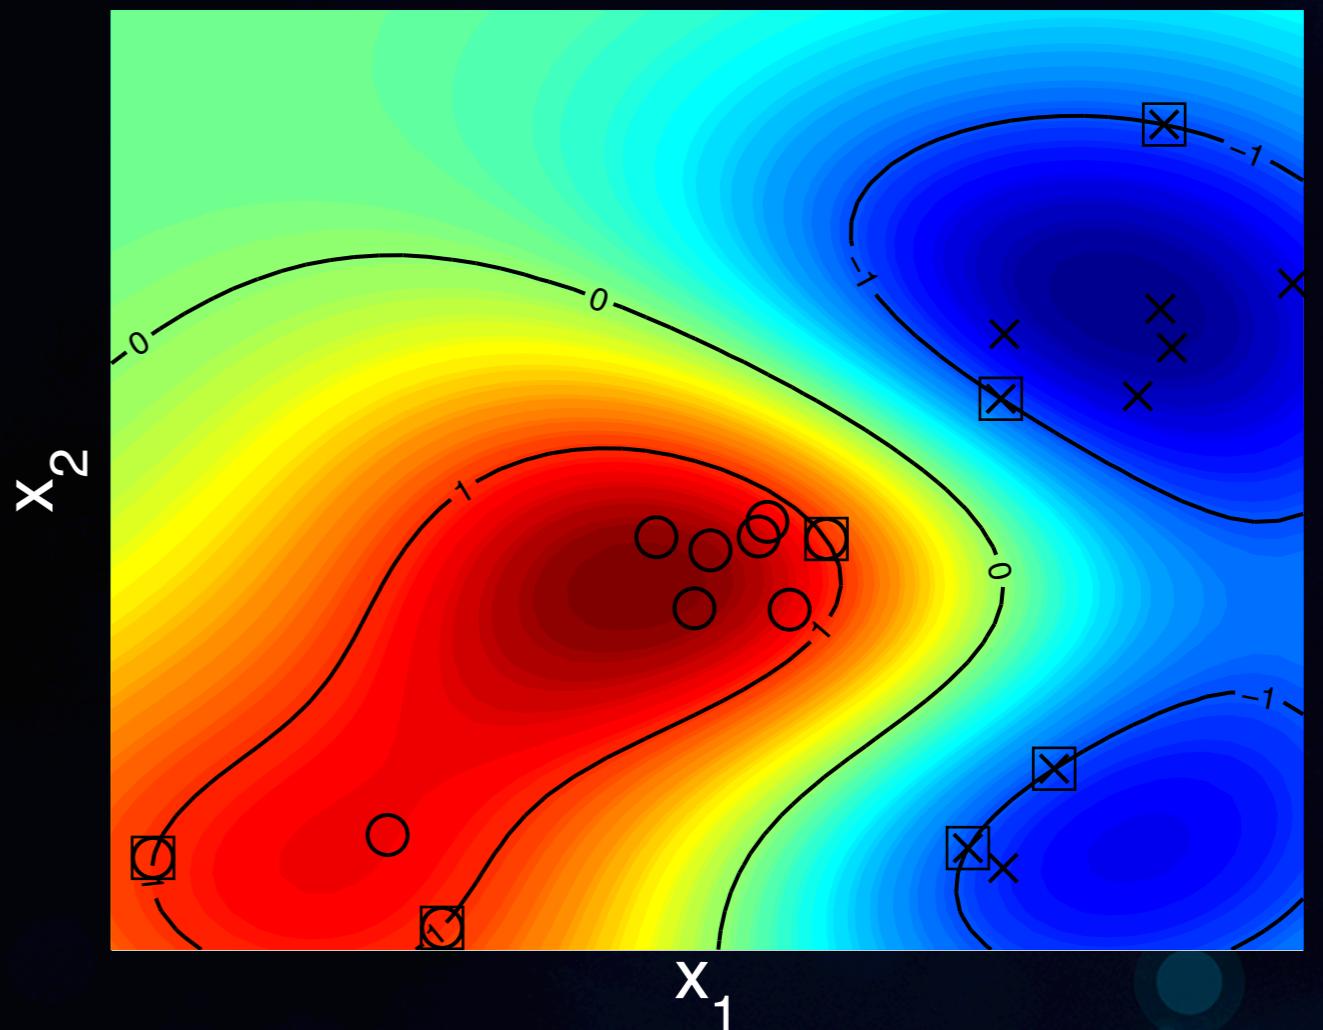
$$\text{if } y_i = -1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b \leq -1$$

$$\text{for } \forall y \\ y_i(\langle \mathbf{w}, \mathbf{x}_i \rangle + b) \geq 1$$

Optimization Function:

$$\text{minimize } \frac{1}{2} \|\mathbf{w}\|^2$$

$$\text{subject to } y_i(\langle \mathbf{w}, \mathbf{x}_i \rangle + b) - 1 \geq 0 \text{ for all } 1 \leq i \leq m$$



# Neural Networks Became Uncool

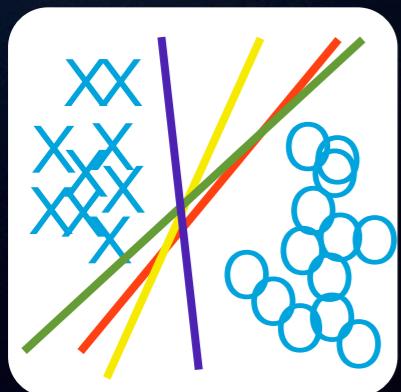
## SVMs

- SVMs killed the credibility of neural network research
- They soon went the way of other 1980's technology
- Neural networks about as uncool as you could possibly get...

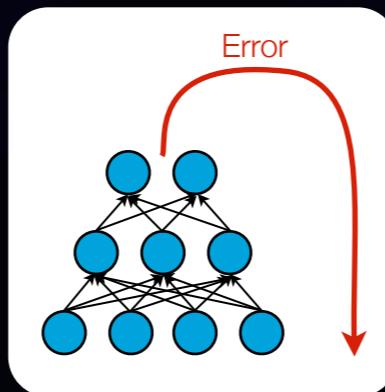


# Deep Learning

A long, long time ago...



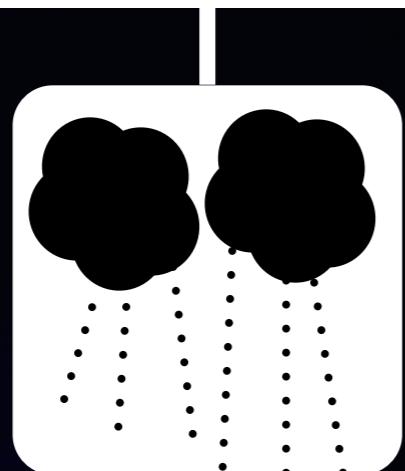
AI Winter  
1969



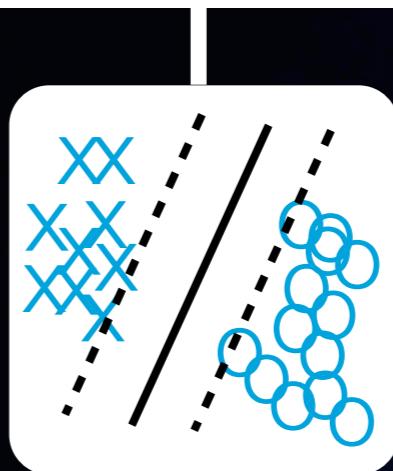
Support Vector  
Machines  
1995



1959  
Perceptrons



1986  
Artificial Neural  
Networks

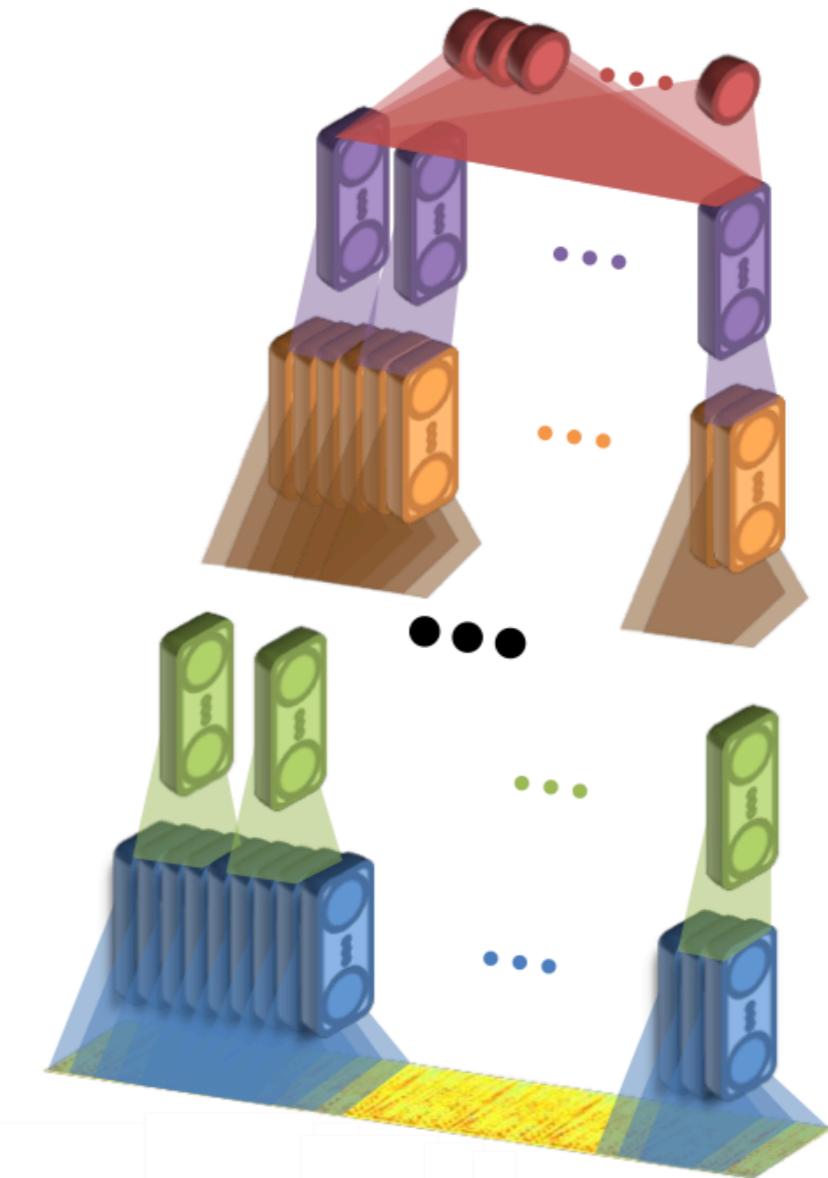


LeNet  
1998

# Convolutional Neural Networks

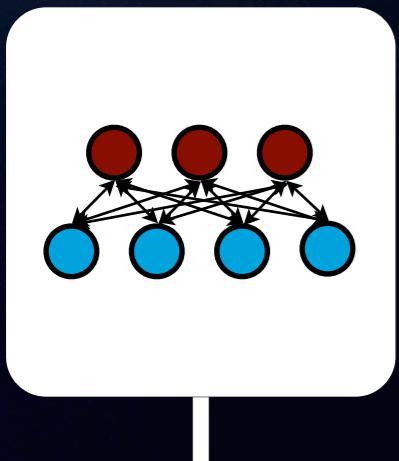
Interesting work did continue

- Initial work focused on image and handwritten digit detection
  - Problems where scale/ rotation invariance are important
- Using convolutional functions allows taking variability of shapes
  - Convolutional layers (linear)
  - Max pooling (non-linear)



# Deep Learning

A little more recently...



2002  
Restricted  
Boltzman  
Machines

# Restricted Boltzman Machines

Building blocks of “deep belief networks”

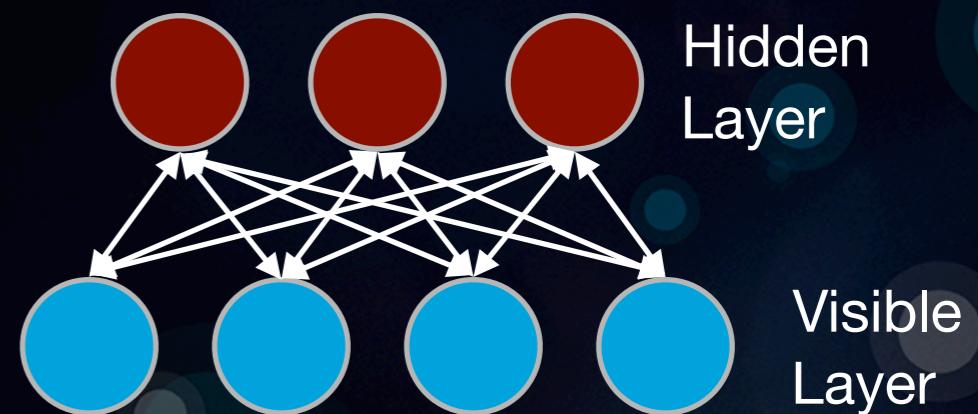
- Restricted Boltzman machines (RBMs) are a type of log-linear Markov random field with hidden units
- Two layer architecture, visible units  $v$  and hidden units  $h$

Energy Function:

$$E(v, h) = - \sum_{i \in \text{freq}} b_i v_i - \sum_{j \in \text{features}} c_j h_j - \sum_{i,j} v_i h_j w_{ij}$$

Model Distribution:

$$P(v) = \sum_h P(v, h) = \sum_h \frac{e^{-E(v, h)}}{Z}$$



# Restricted Boltzman Machines

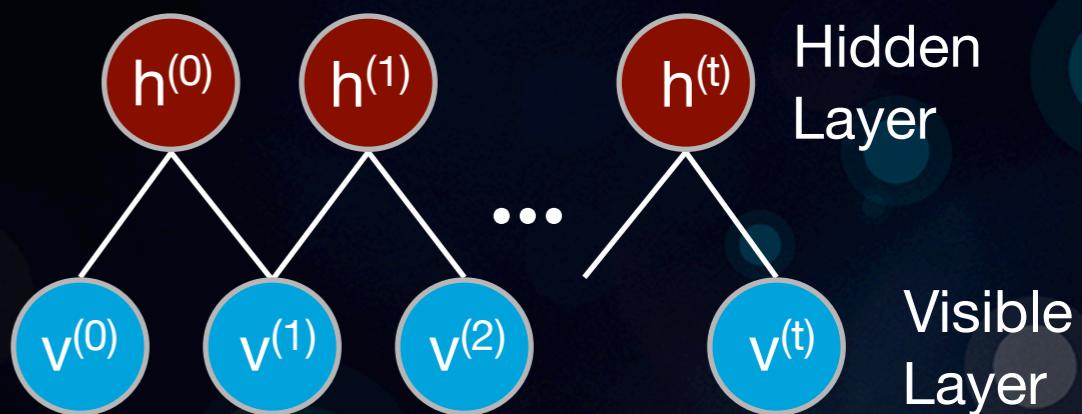
## Inference

- In order to optimize this function it is necessary to first determine the values of the hidden variables
- Gibbs sampling is used in this case
- In theory it should be run until the chain converges (i.e., values no longer change)

$$\mathbf{h}^{(n+1)} \sim \sigma(b_j + \sum_i w_{ij} v_i),$$

$$\mathbf{v}^{(n+1)} \sim \sigma(a_i + \sum_j w_{ij} h_j)$$

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$



# Restricted Boltzman Machines

## Learning

- In practice, the general approach is to run the chain for just 1 step
- Update the parameter estimates for the model
- Weighted by a learning rate  $\epsilon$
- This is done iteratively for some number of epochs

$$\frac{\partial \log P(\mathbf{v})}{\partial W} = \mathbf{v}' E[\mathbf{h}|\mathbf{v}] - E[\mathbf{v}'\mathbf{h}],$$

$$\frac{\partial \log P(\mathbf{v})}{\partial a} = E[\mathbf{h}|\mathbf{v}] - E[\mathbf{h}],$$

$$\frac{\partial \log P(\mathbf{v})}{\partial b} = \mathbf{v} - E[\mathbf{v}].$$

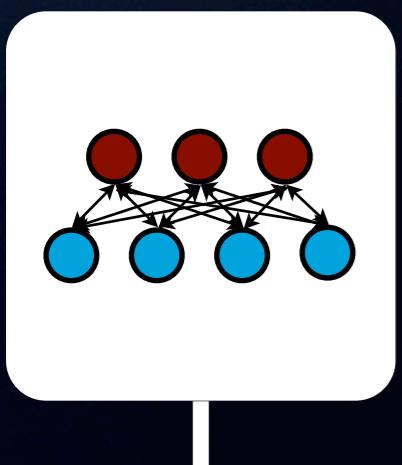
$$W = W + \epsilon(\mathbf{v}^{(0)}\hat{\mathbf{h}}^{(0)'} - \mathbf{v}^{(1)}\hat{\mathbf{h}}^{(1)'}),$$

$$a = a + \epsilon(\hat{\mathbf{h}}^{(0)} - \hat{\mathbf{h}}^{(2)}),$$

$$b = b + \epsilon(\mathbf{v}^{(0)} - \mathbf{v}^{(1)})$$

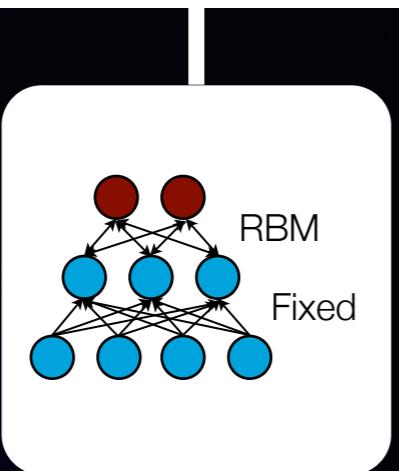
# Deep Learning

A little more recently...



Greedy-Wise  
Pretraining  
2006

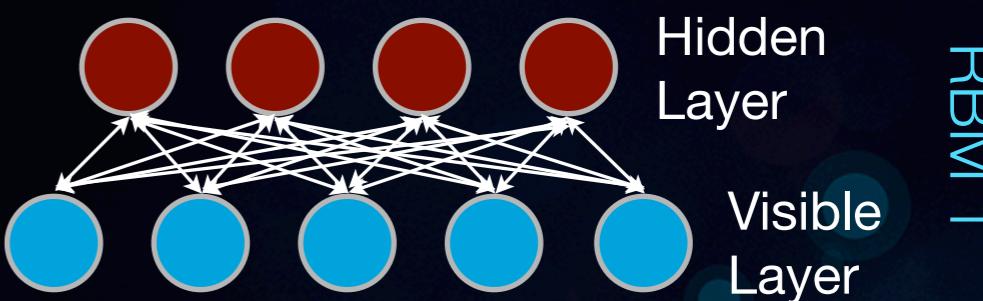
2002  
Restricted  
Boltzman  
Machines



# Learning Deep Autoencoders

## Stacked RBMs

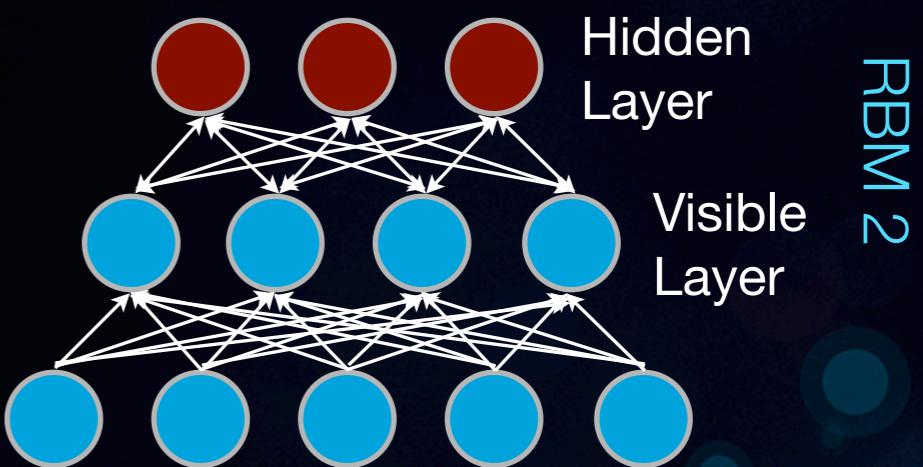
- Stacked RBMs shown as a powerful way to initialize autoencoders
- First we train RBMs one layer at a time
  - Previous hidden layer becomes current visible



# Learning Deep Autoencoders

## Stacked RBMs

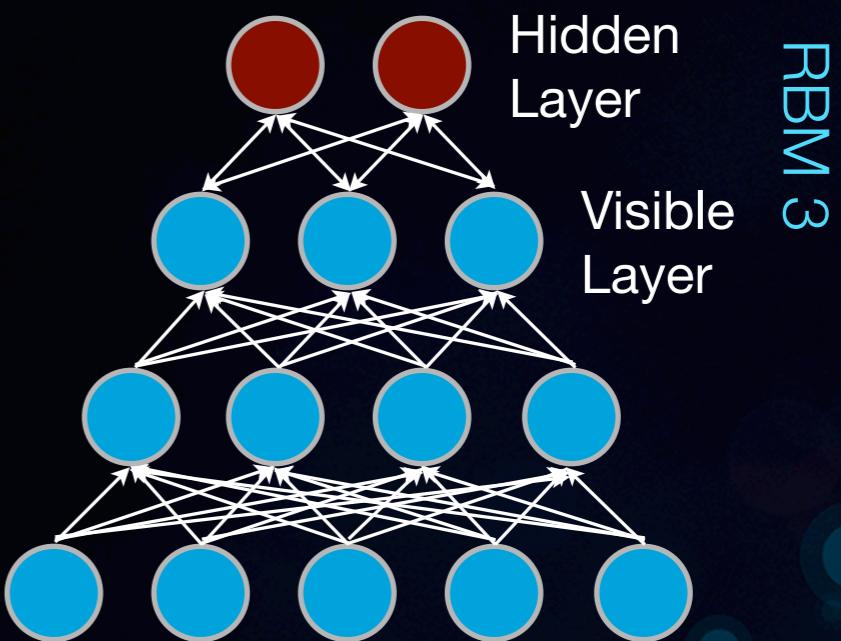
- Stacked RBMs shown as a powerful way to initialize autoencoders
- First we train RBMs one layer at a time
  - Previous hidden layer becomes current visible



# Learning Deep Autoencoders

## Stacked RBMs

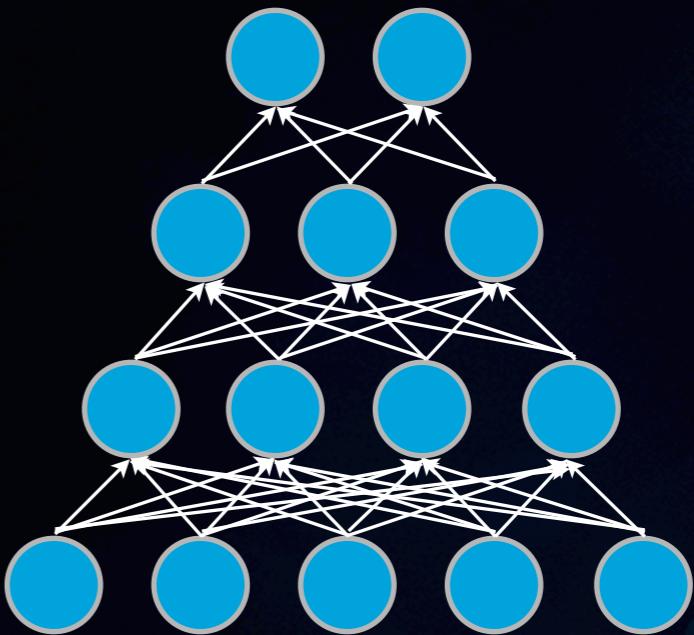
- Stacked RBMs shown as a powerful way to initialize autoencoders
- First we train RBMs one layer at a time
  - Previous hidden layer becomes current visible



# Learning Deep Autoencoders

## Stacked RBMs

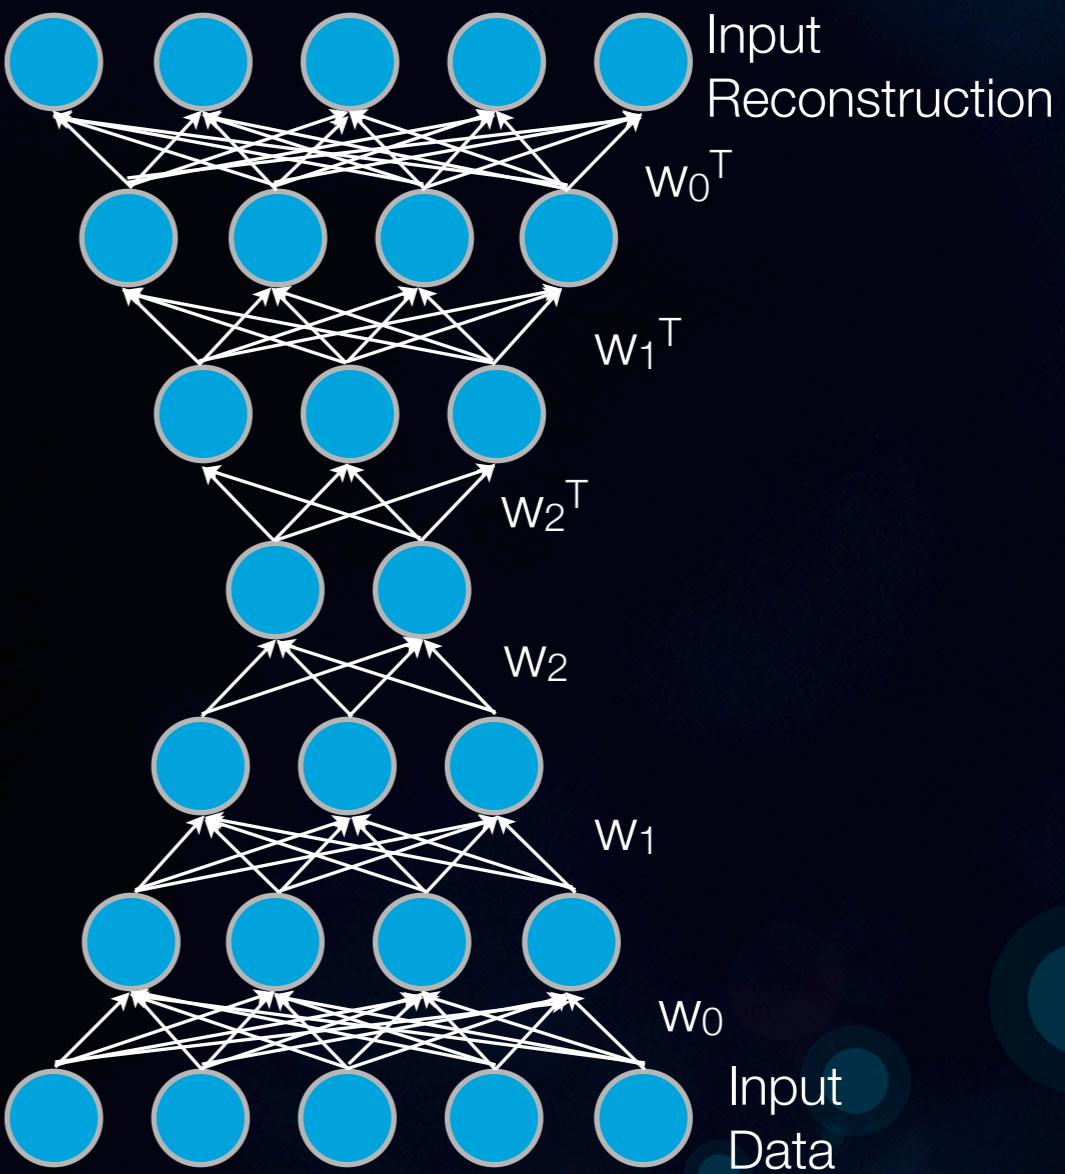
- Stacked RBMs shown as a powerful way to initialize autoencoders
- First we train RBMs one layer at a time
  - Previous hidden layer becomes current visible



# Learning Deep Autoencoders

## Stacked RBMs

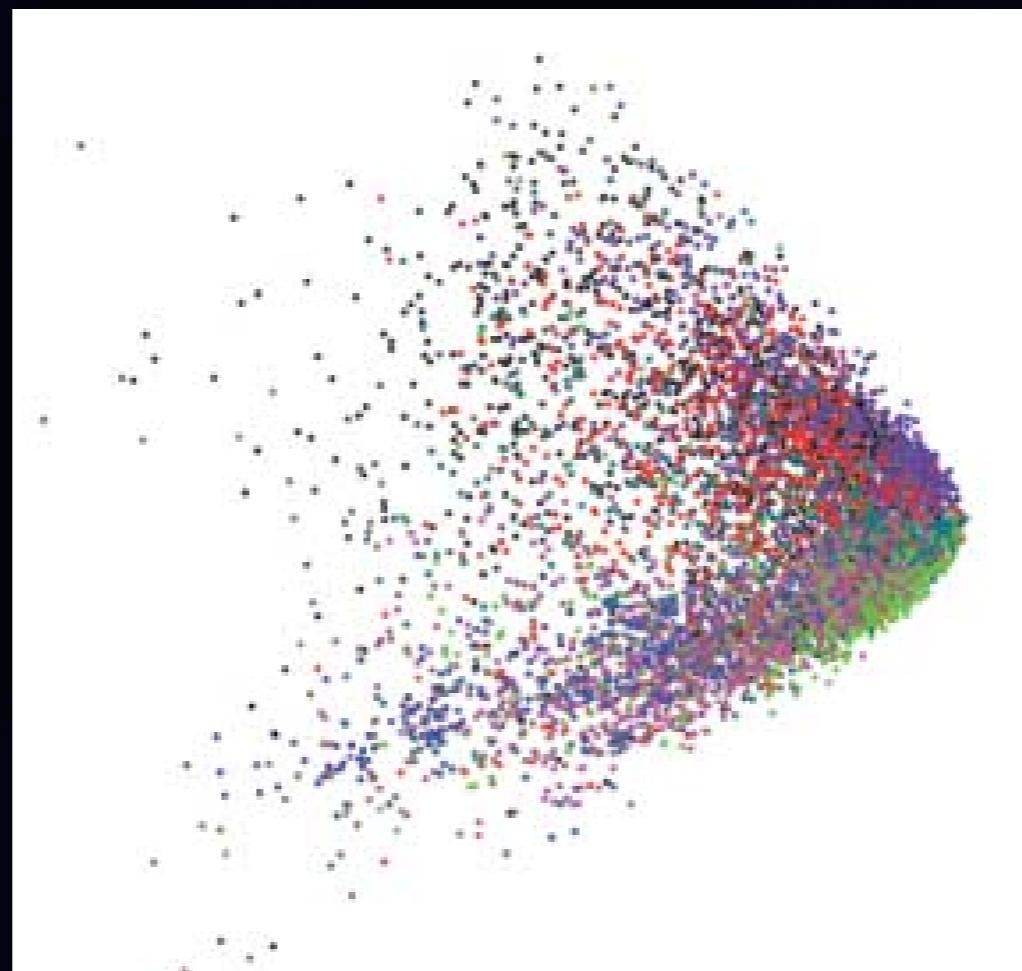
- Stacked RBMs shown as a powerful way to initialize autoencoders
- First we train RBMs one layer at a time
  - Previous hidden layer becomes current visible



# Learning Deep Autoencoders

## Stacking RBMs

Two dimensional latent semantic analysis (Figure 4, B)

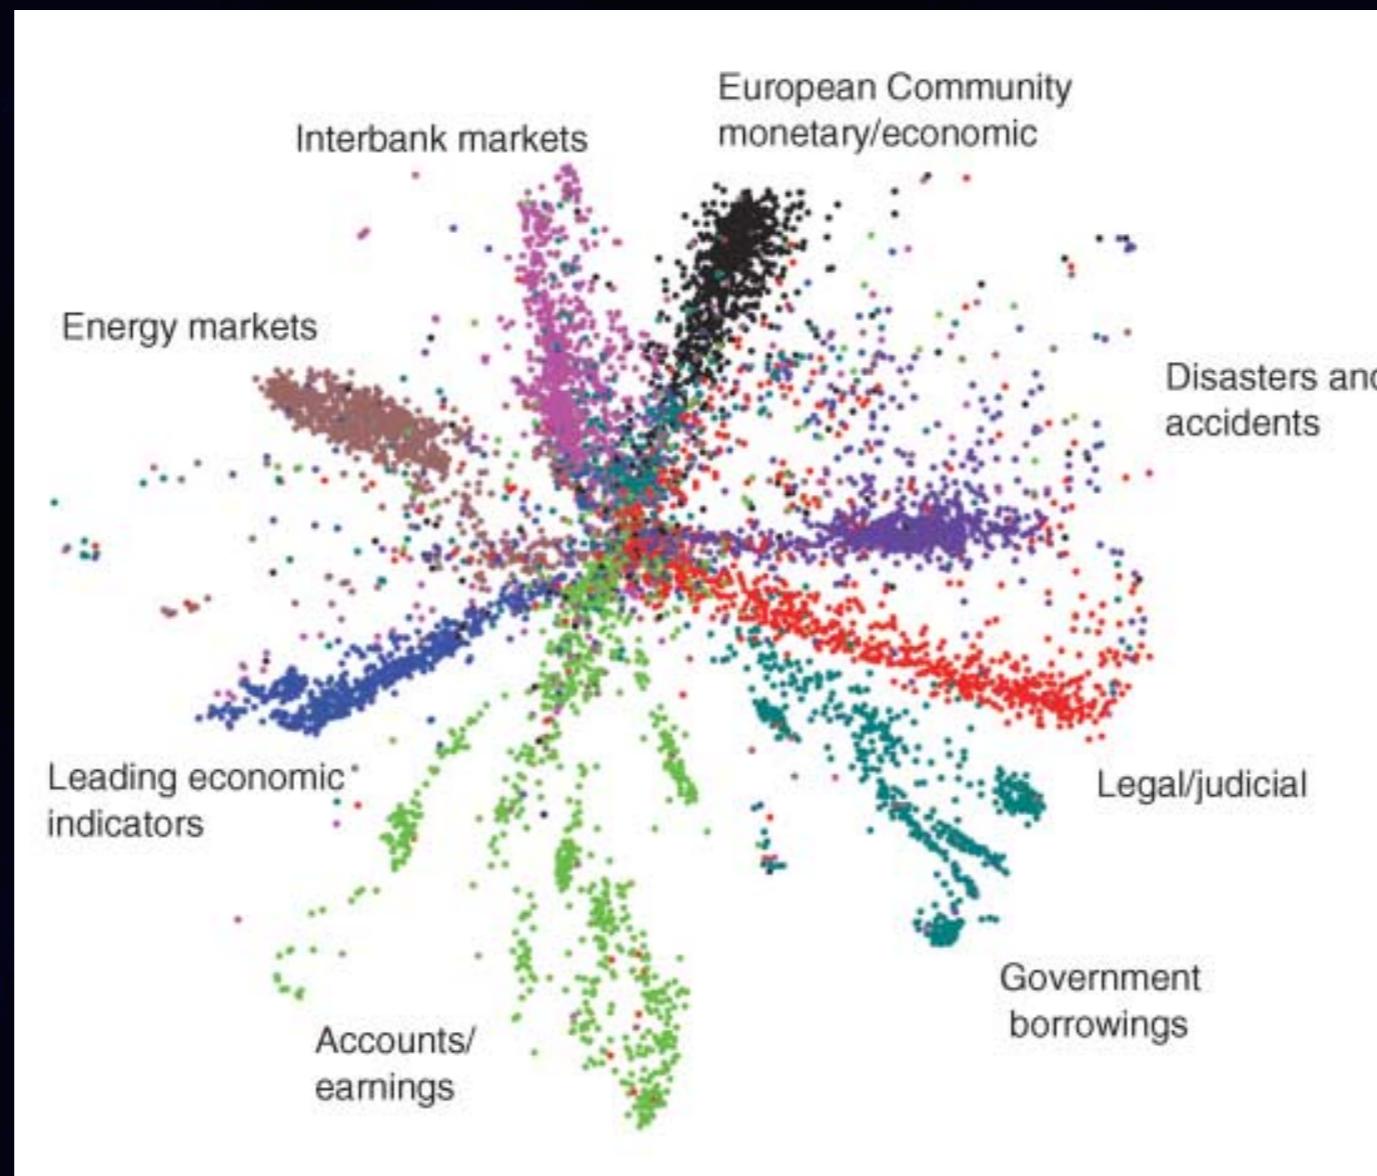


G. E. Hinton and R. R. Salakhutdinov (2006). Reducing the Dimensionality of Data with Neural Networks. *Science*.

# Learning Deep Autoencoders

## Stacking RBMs

2000-500-250-125-2 Autoencoder (Figure 4, C)



G. E. Hinton and R. R. Salakhutdinov (2006). Reducing the Dimensionality of Data with Neural Networks. *Science*.

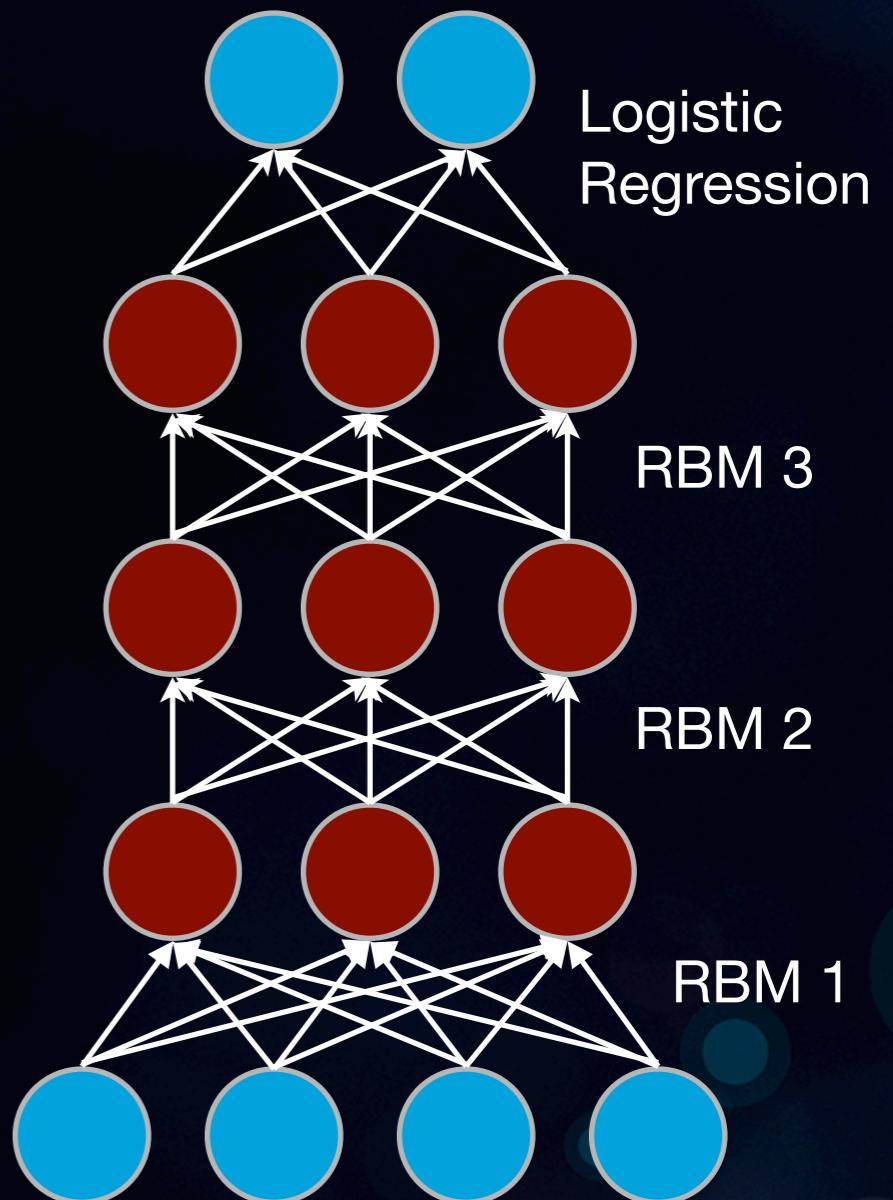
# Greedy-Wise Pre-training

## Initializing model with RBMs

- DBNs are pre-trained by “greedily” stacking restricted Boltzmann machines (RBMs)
- Logistic regression is attached last
- Model is fine-tuned with gradient descent:

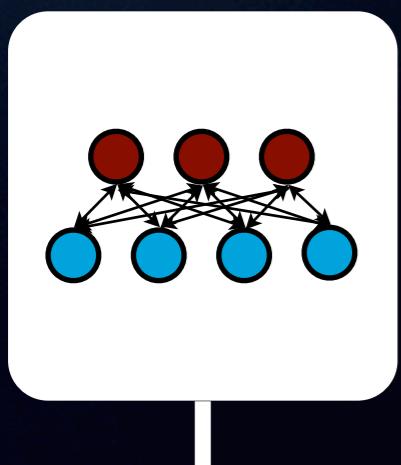
$$\theta^{(t+1)} = \theta^{(t)} - \epsilon \Delta E(\theta^{(t)})$$

- Model parameters  $\theta$
- Learning rate  $\epsilon$  (chosen experimentally)
- Error function negative log likelihood

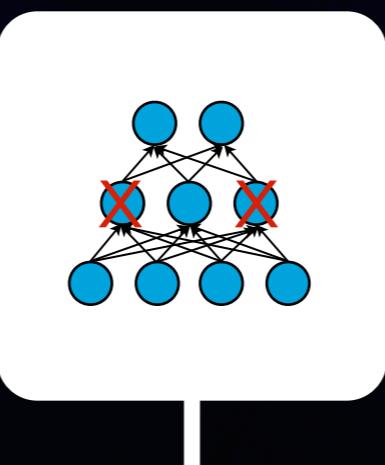


# Deep Learning

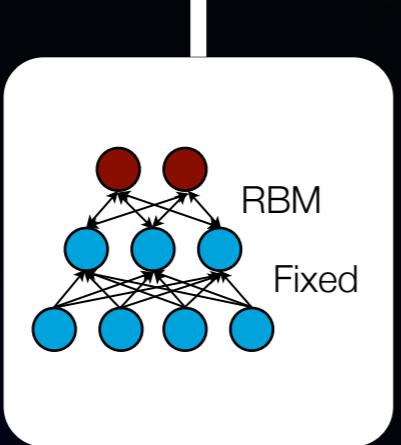
A little more recently...



Greedy-Wise  
Pretraining  
2006



2002  
Restricted  
Boltzman  
Machines



2012  
Dropout

# Dropout

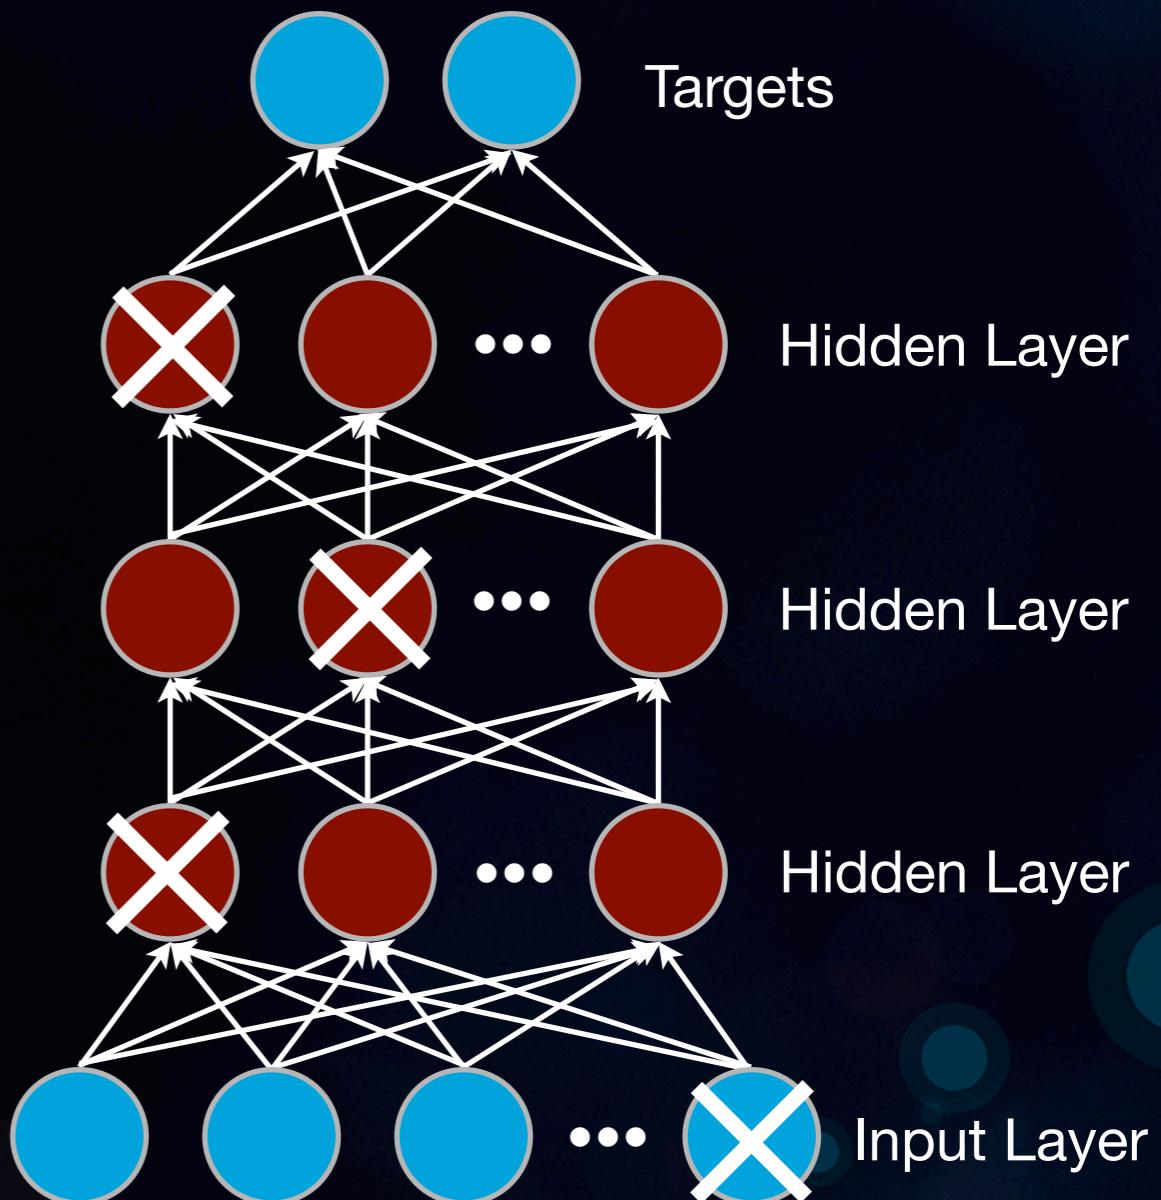
## Regularization

- So far the methods discussed use very simplistic regularization strategies
  - A validation set is held out (separate from training and test data)
  - On each iteration we verify that error is reducing on both the training and on validation set
  - Since the validation set is unseen by the algorithm, this will help to prevent over fitting

# Dropout

## Regularization

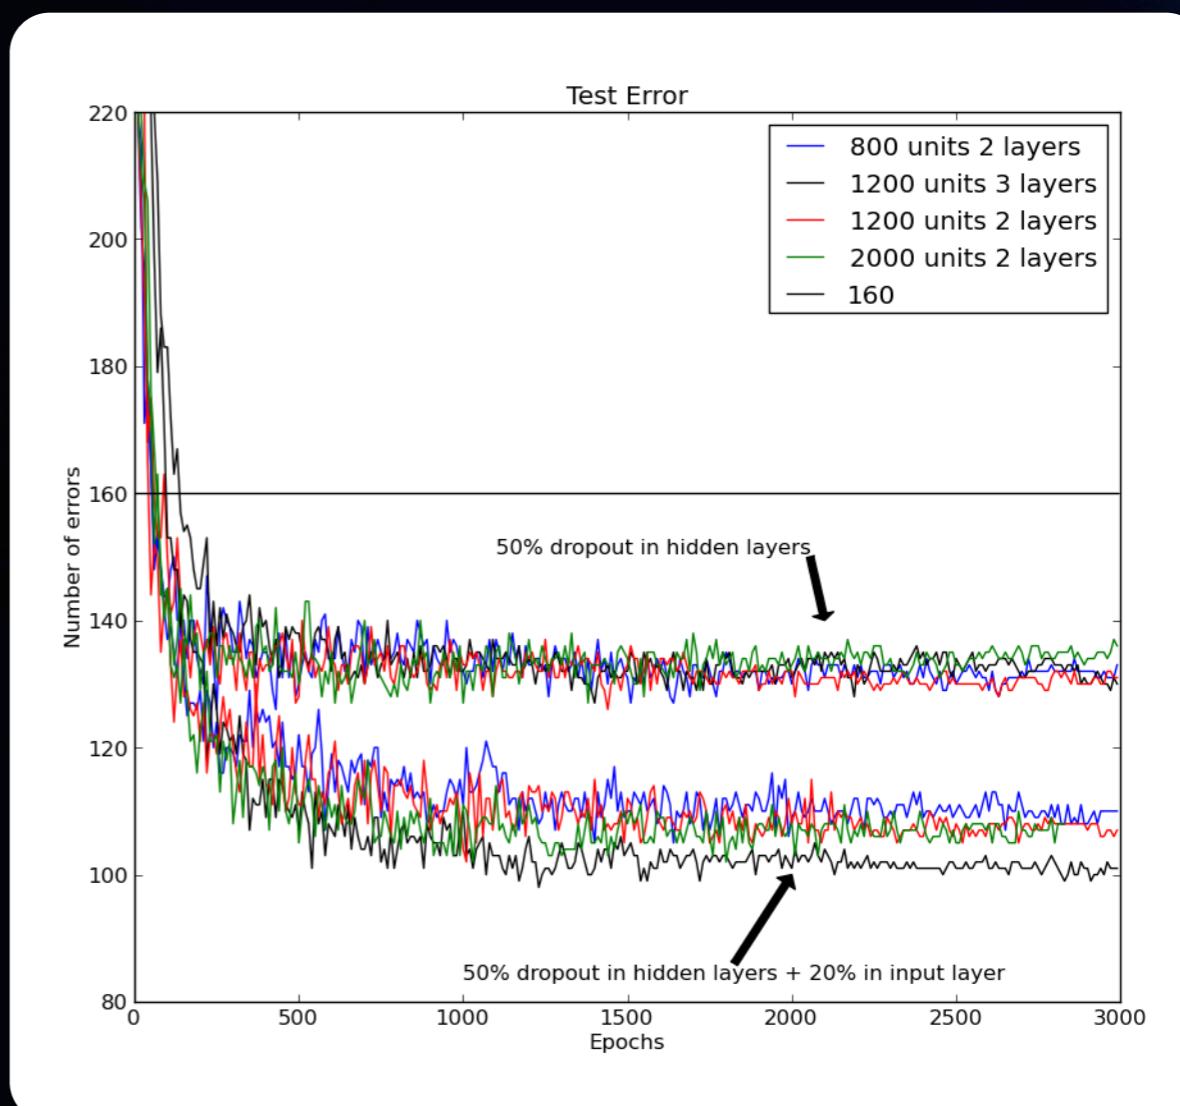
- Dropout takes things a step further
- Independently set each hidden unit's activity to 0 with 0.5 probability
- Correlations between nodes cannot develop
- Also called co-adaptions



# Dropout

## MNIST

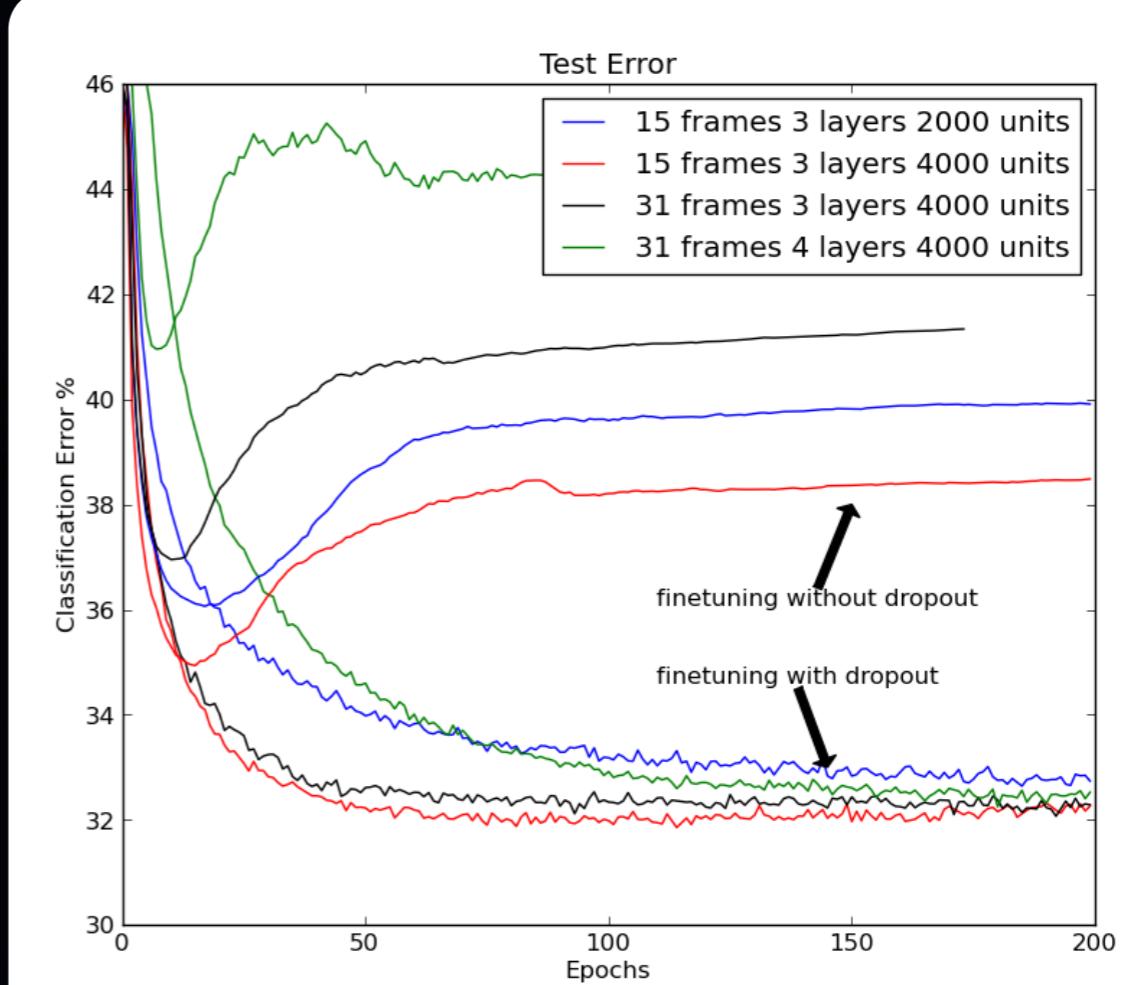
- Handwritten digit recognition
- 60,000 20x20 training images
- 20,000 test images
- Previously best result was 160 errors (black line)



# Dropout

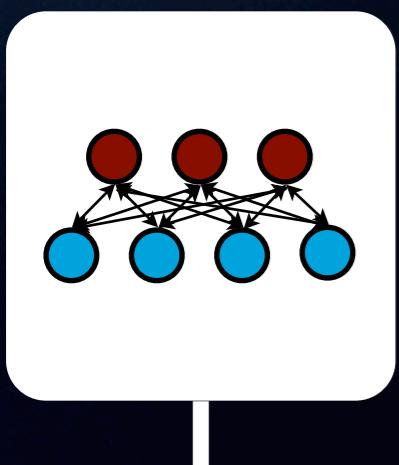
## TIMIT

- Recognition of clean speech (no noise)
- Pre-trained DBNs had been shown previously to outperform GMMs
- Used in combination with hidden Markov model

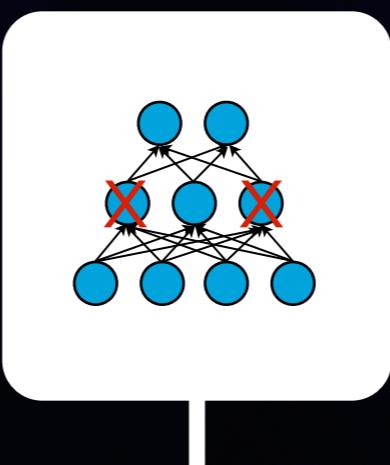


# Deep Learning

A little more recently...

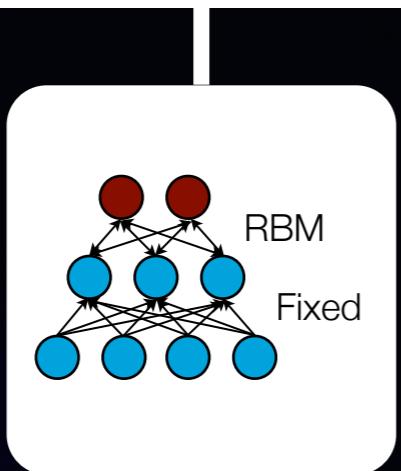


Greedy-Wise  
Pretraining  
2006

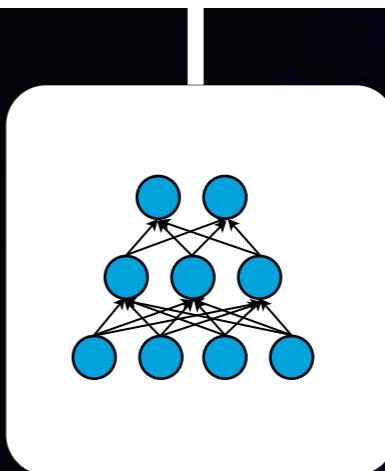


Artificial Neural  
Networks  
1986/2013

2002  
Restricted  
Boltzman  
Machines



2012  
Dropout



# Artificial Neural Networks

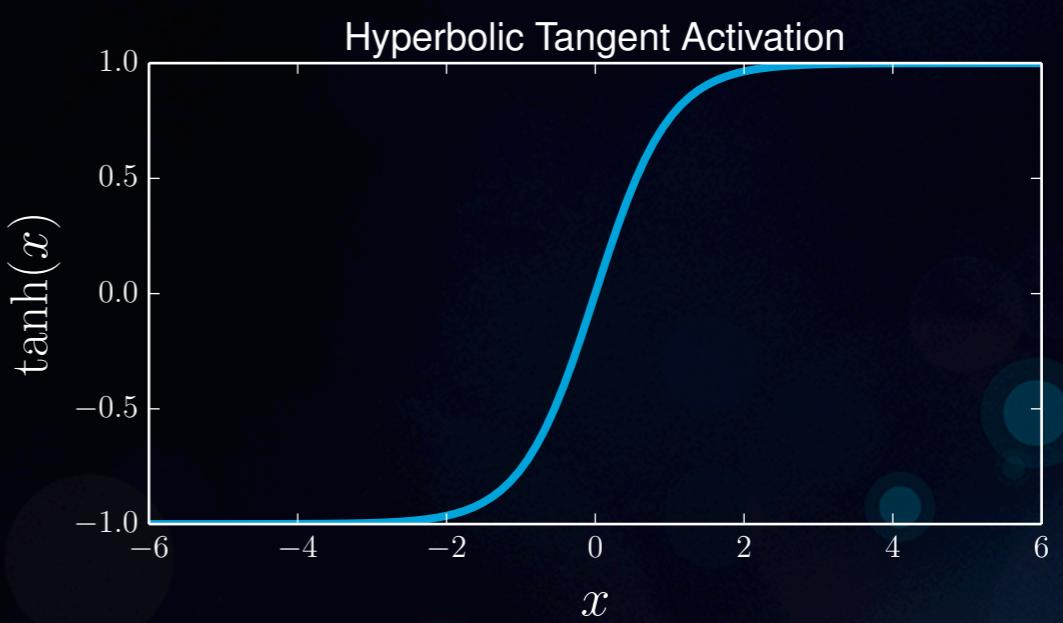
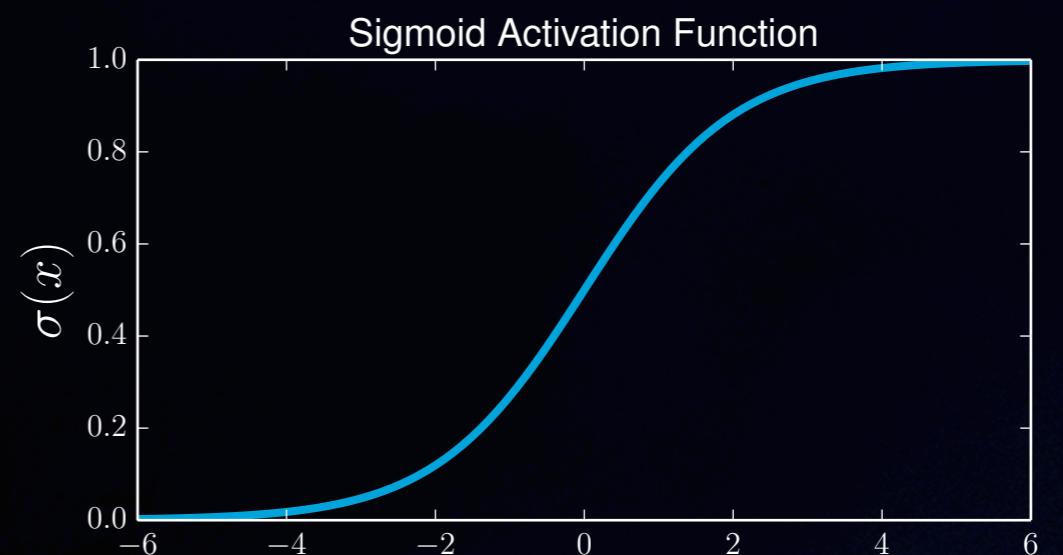
Back to 1986!

- Many researchers are now claiming that neither pre-training nor dropout are necessary given enough training data
- These methods are nearly identical to the original neural networks of 1986
- The only major change is the use of rectified linear units
  - And **\*a lot\*** more data
  - As Geoff Hinton likes to put it, we just had to wait long enough for our datasets and computing power to scale

# Rectified Linear Units

Change in activation

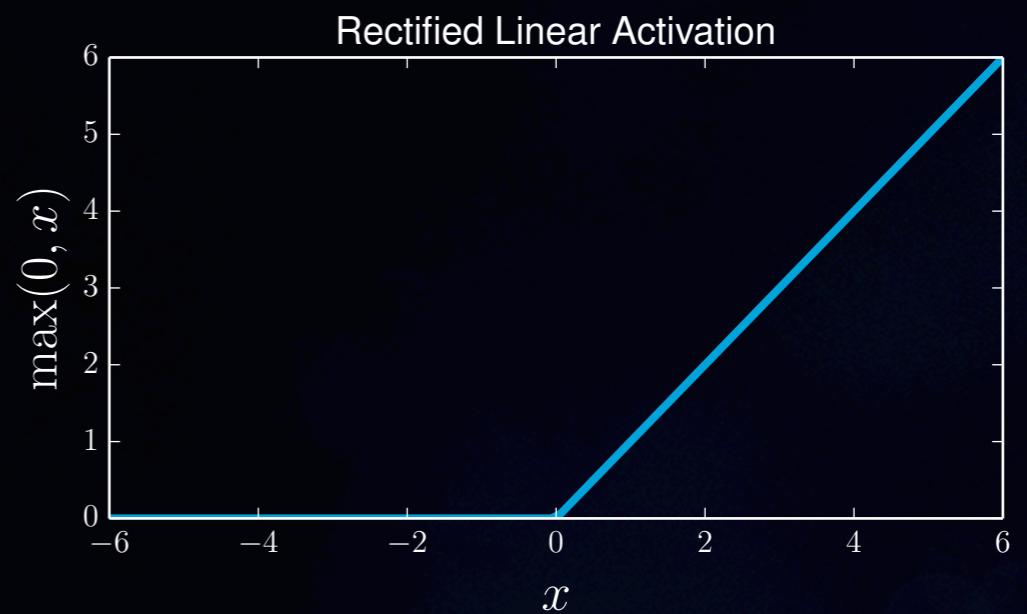
- Computing non-linear activations is computationally expensive
- Each dimension of each example must be evaluated
- With rectified linear units the values remain unchanged (or set to 0)



# Rectified Linear Units

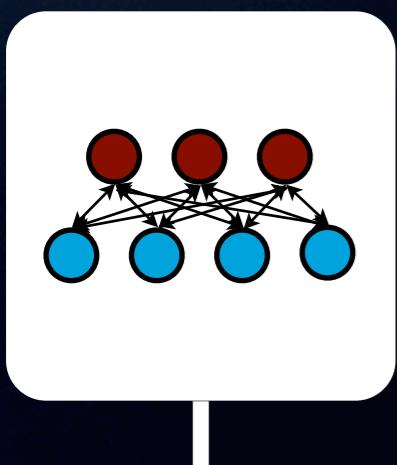
Change in activation

- Computing non-linear activations is computationally expensive
- Each dimension of each example must be evaluated
- With rectified linear units the values remain unchanged (or set to 0)

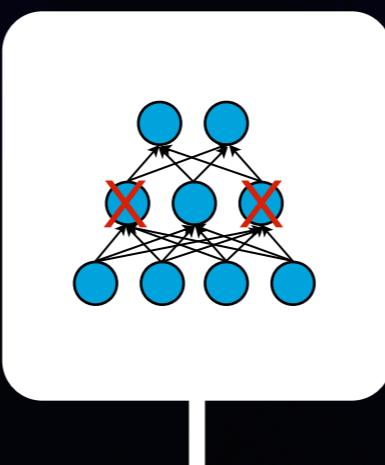


# Deep Learning

A little more recently...



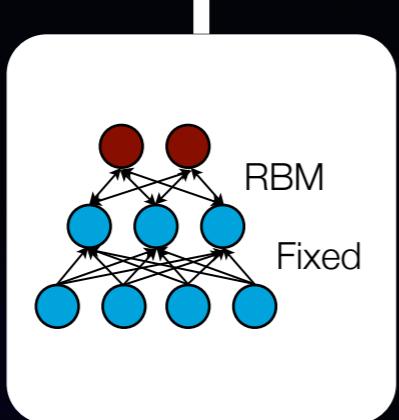
Greedy-Wise  
Pretraining  
2006



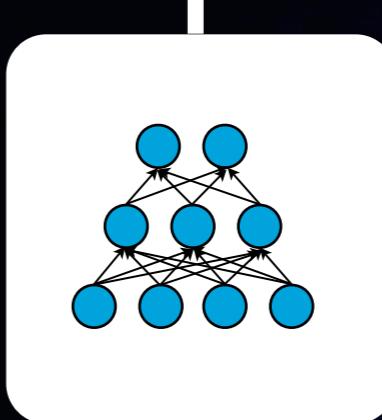
Artificial Neural  
Networks  
1986/2013



2002  
Restricted  
Boltzman  
Machines



2012  
Dropout



The Future

# Music Emotion Recognition

# Music Emotion Recognition

## Overview

### Goal:

- Content-based modeling and prediction of musical emotion
- Applications in sorting, recommendation, playlist generation

### Challenges:

1. Emotional descriptions often lack a singular, well-defined answer making ground truth labeling difficult
2. No dominant acoustic feature has yet arisen for musical emotion
3. Emotions change over time

# Music Emotion Recognition

## Overview

### Goal:

- Content-based modeling and prediction of musical emotion
- Applications in sorting, recommendation, playlist generation

### Challenges:

1. Emotional descriptions often lack a singular, well-defined answer making ground truth labeling difficult
2. No dominant acoustic feature has yet arisen for musical emotion
3. Emotions change over time

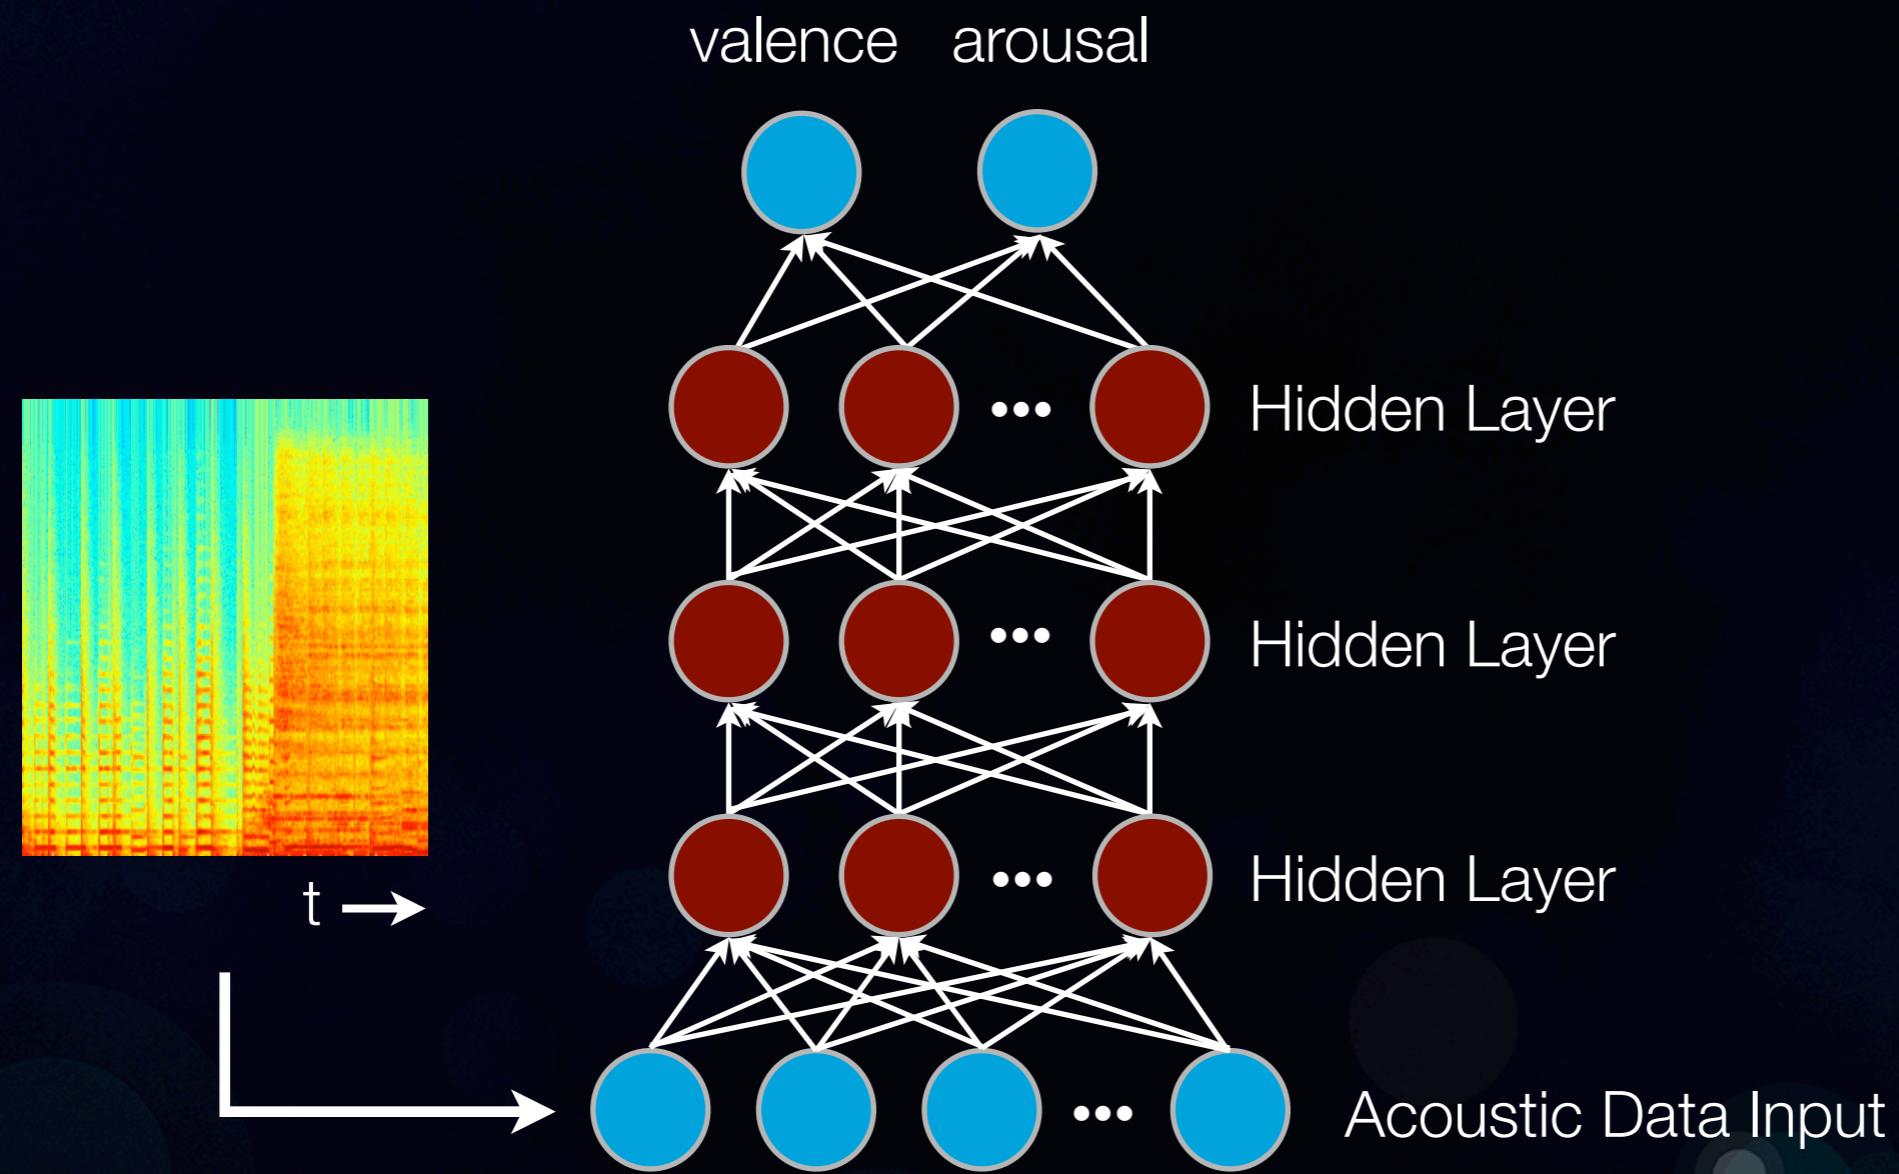
# Ground Truth Data Collection

## MoodSwings Turk

- Re-annotation of MoodSwings dataset with Amazon's Mechanical Turk
- Highly correlated to MoodSwings labels
  - Finds arousal  $r=0.712$ , valence  $r=0.846$
- Annotated with  $16.93 \pm 2.690$  ratings per clip
- Dataset has been made publicly available

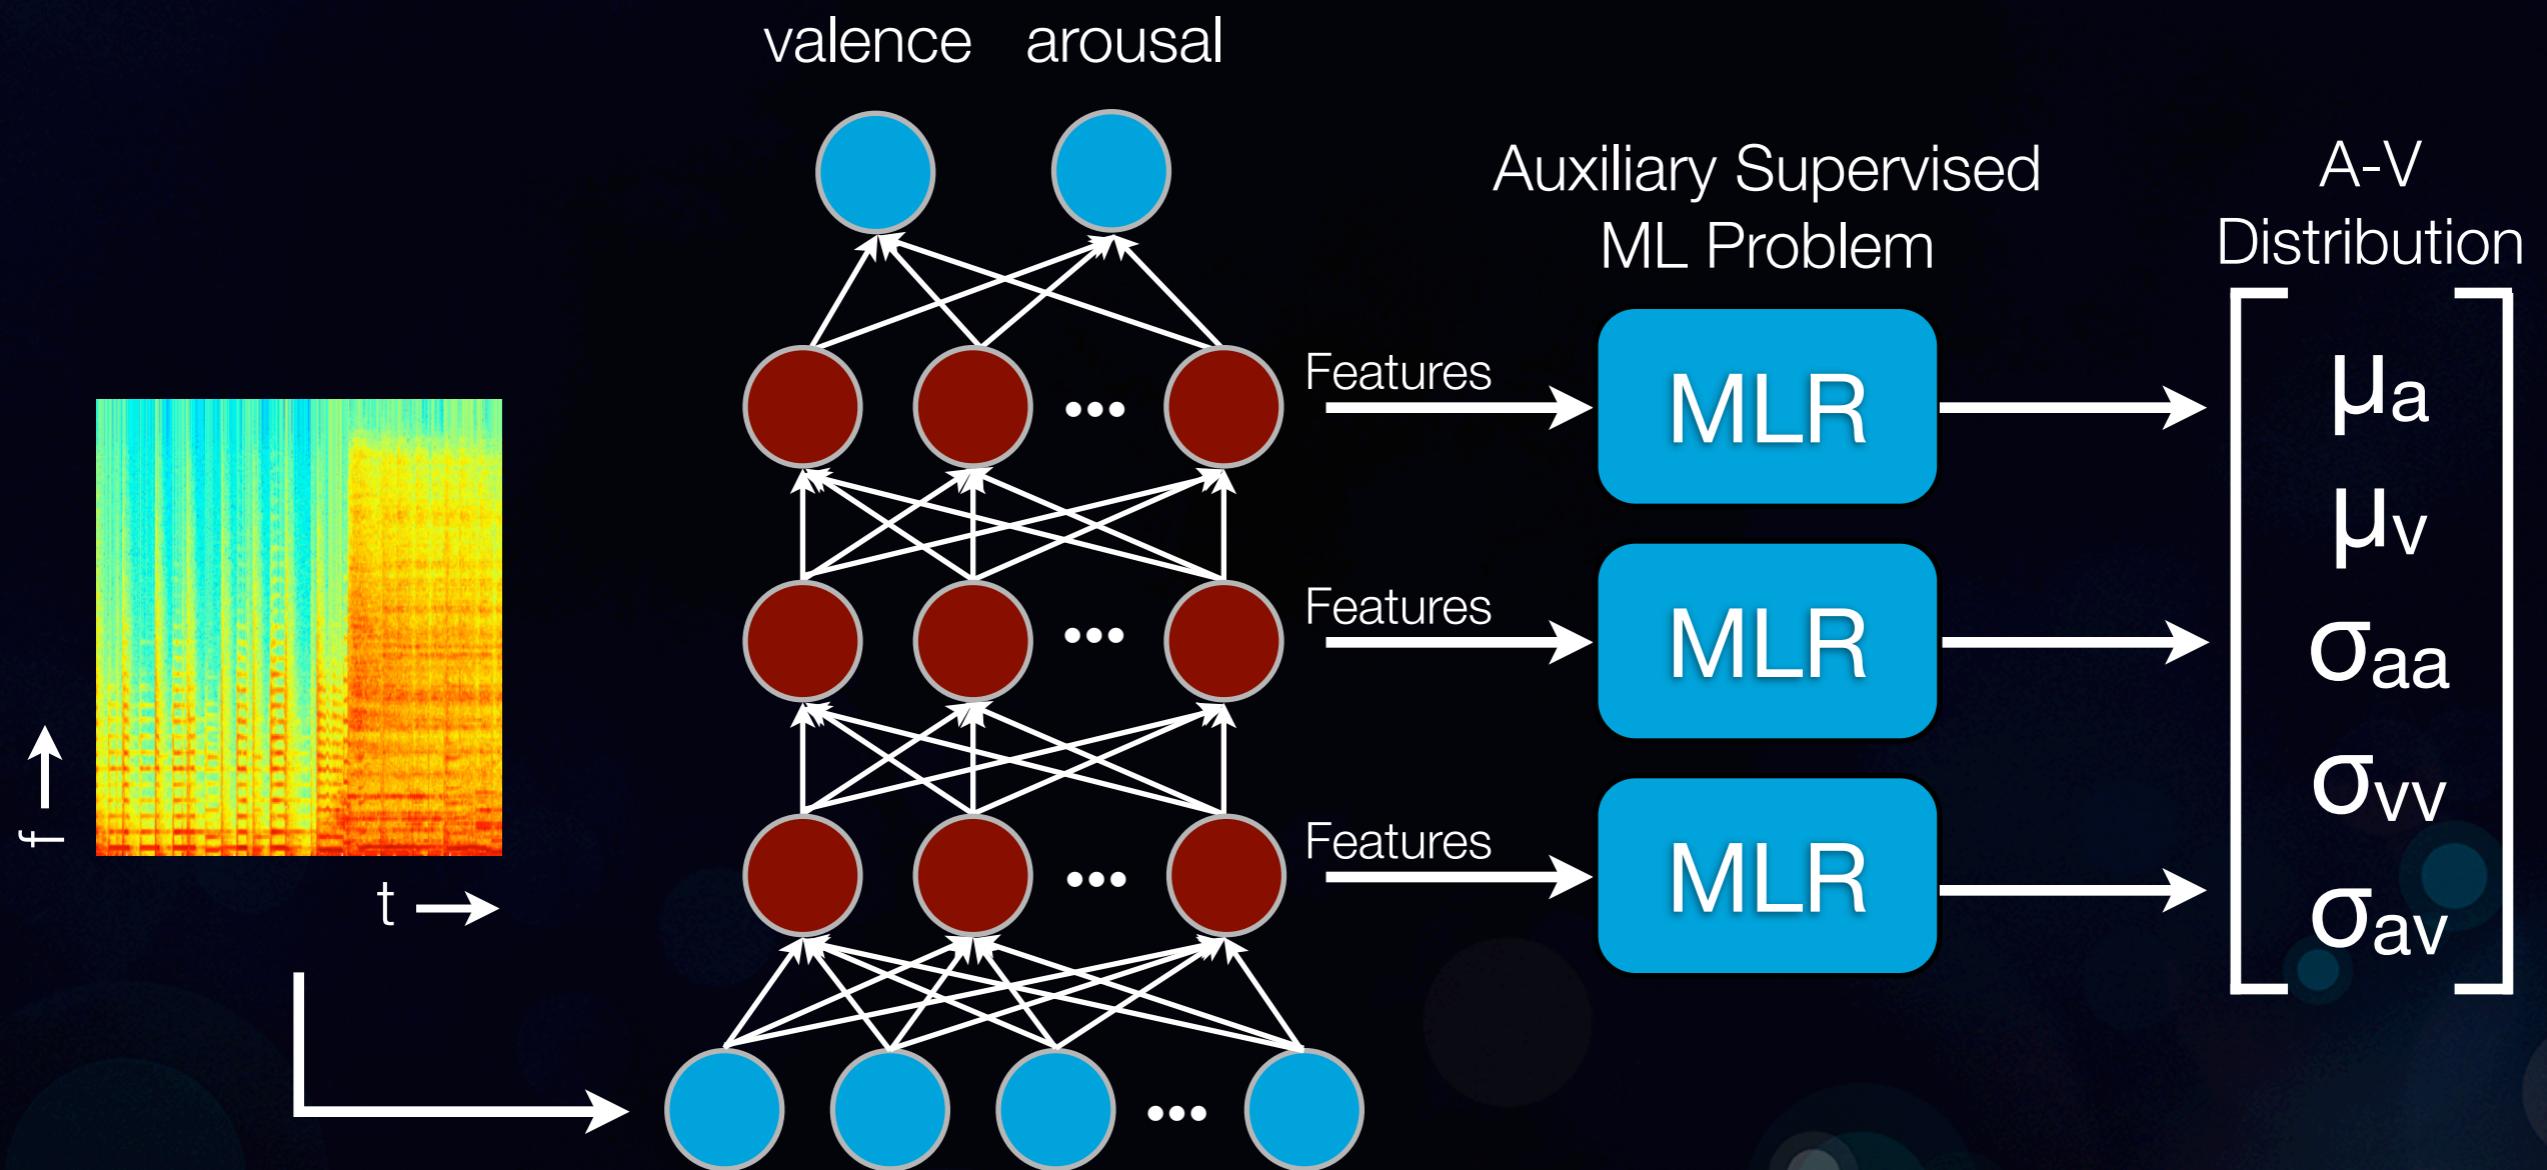
# Framework

Deep learning in music emotion recognition



# Framework

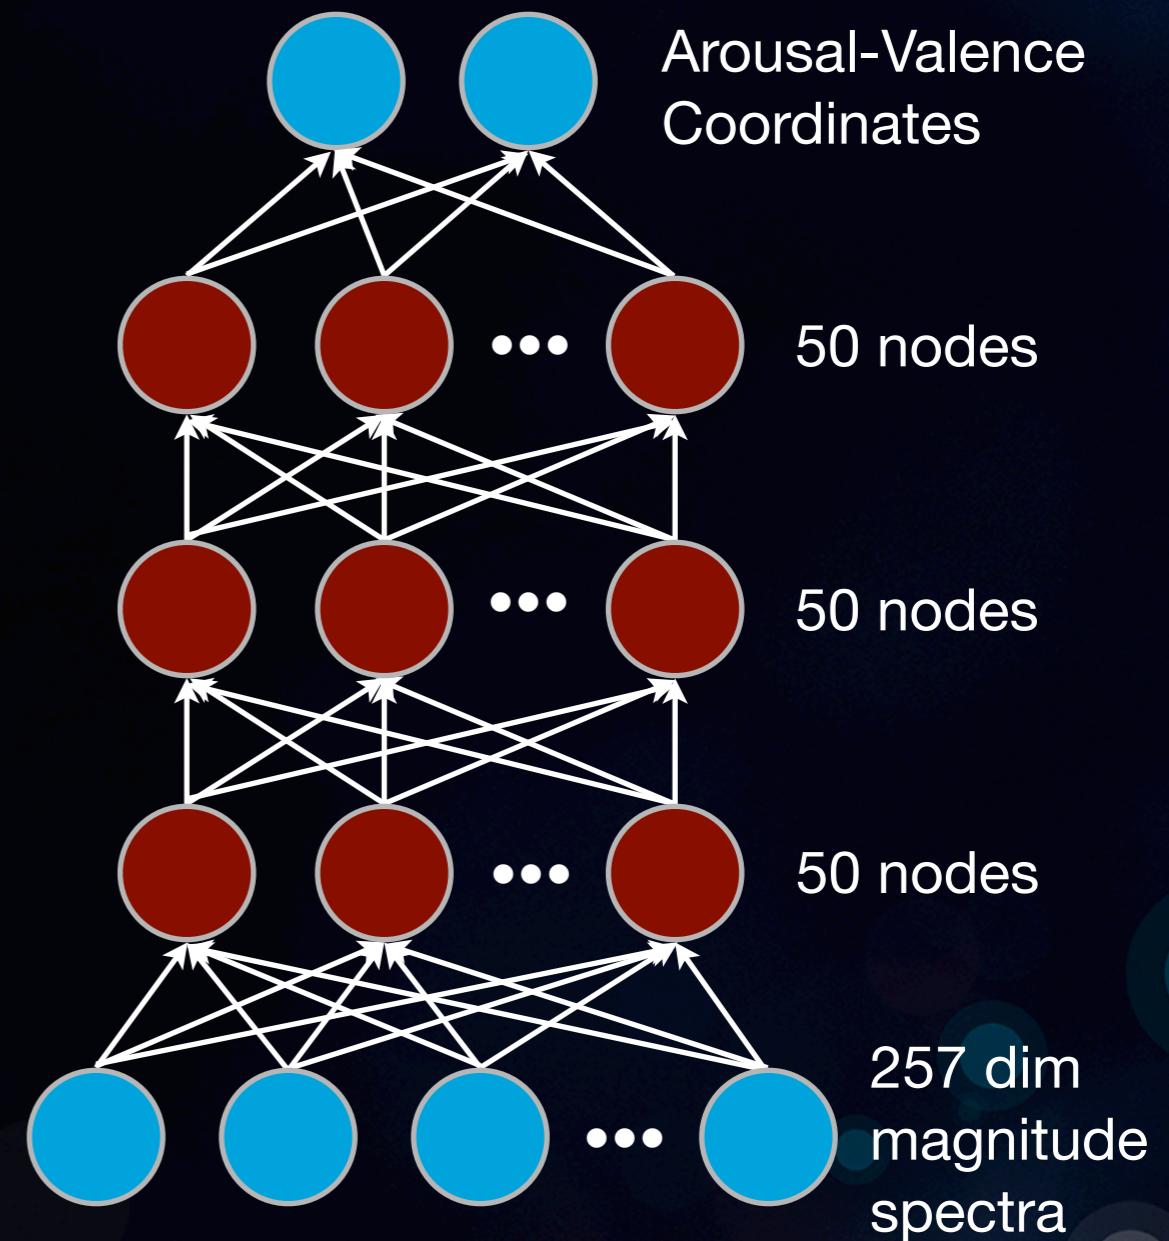
# Deep learning in music emotion recognition



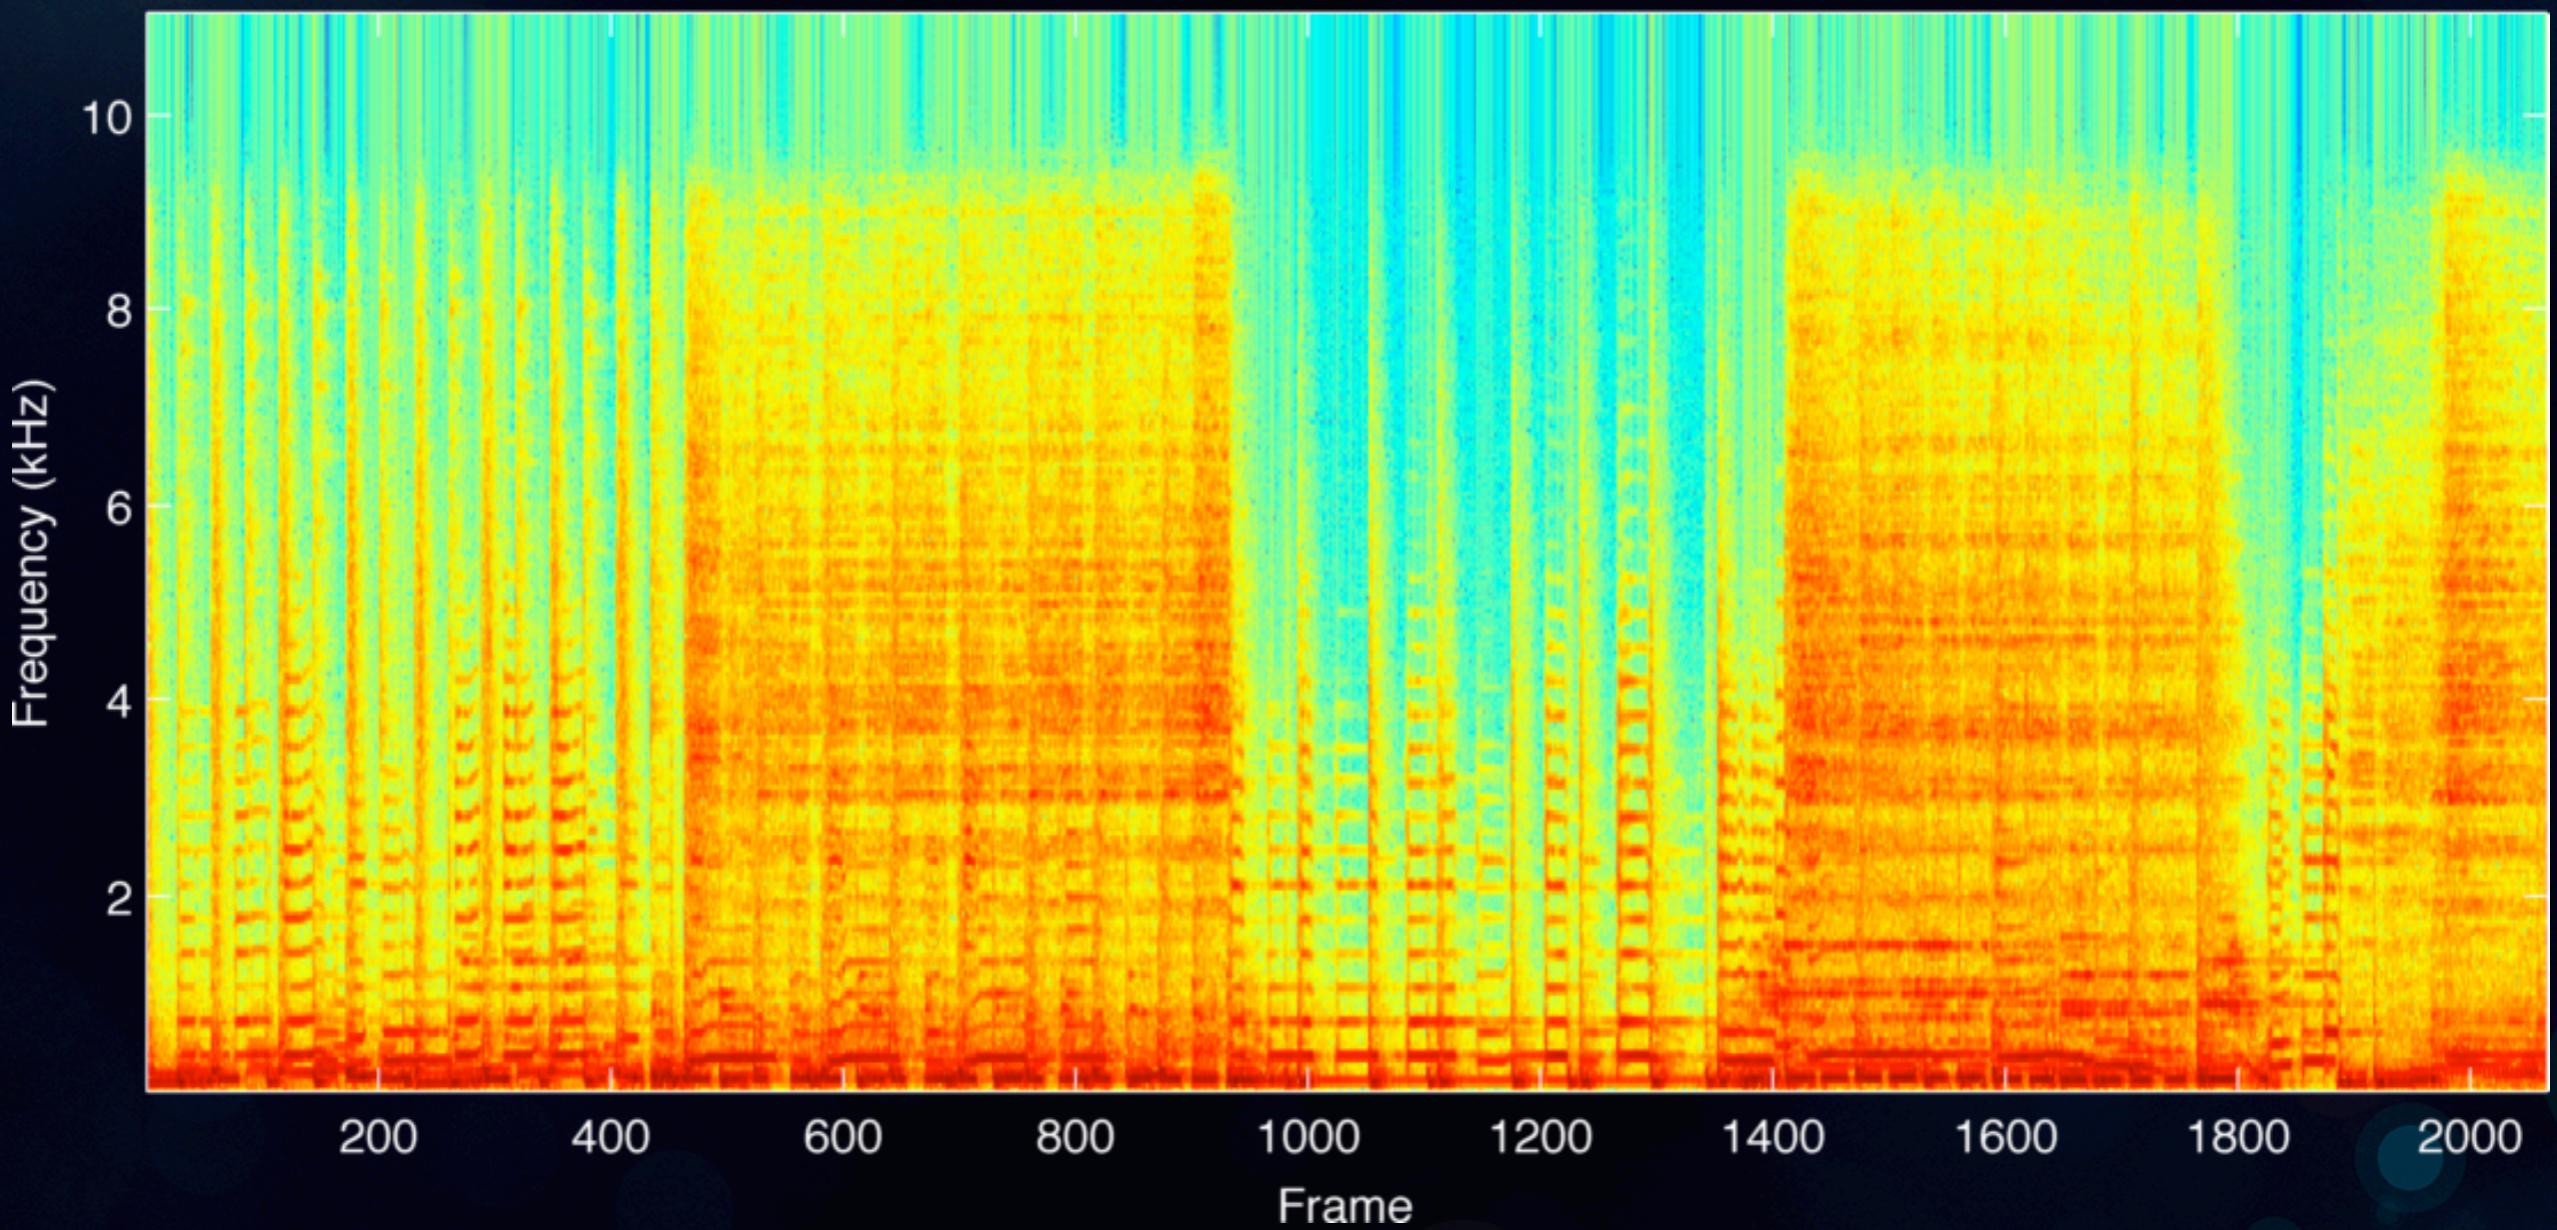
# Feature Learning

## Single STFT Frame

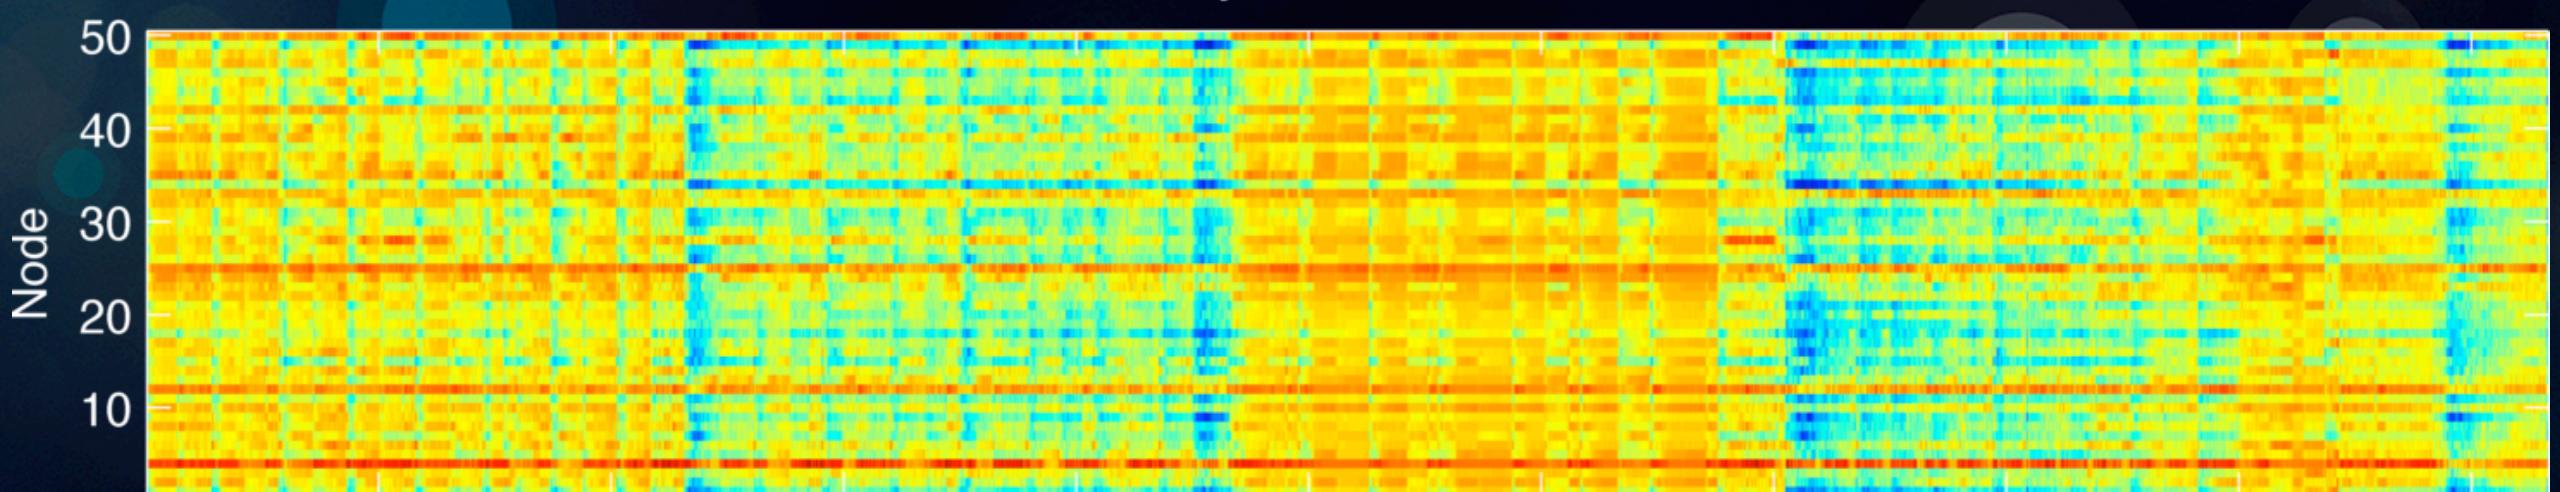
- 50 dimensions are used in the hidden layers (chosen experimentally)
- Model is trained to predict A-V means
- Input is raw magnitude spectra
- ~20msec windows



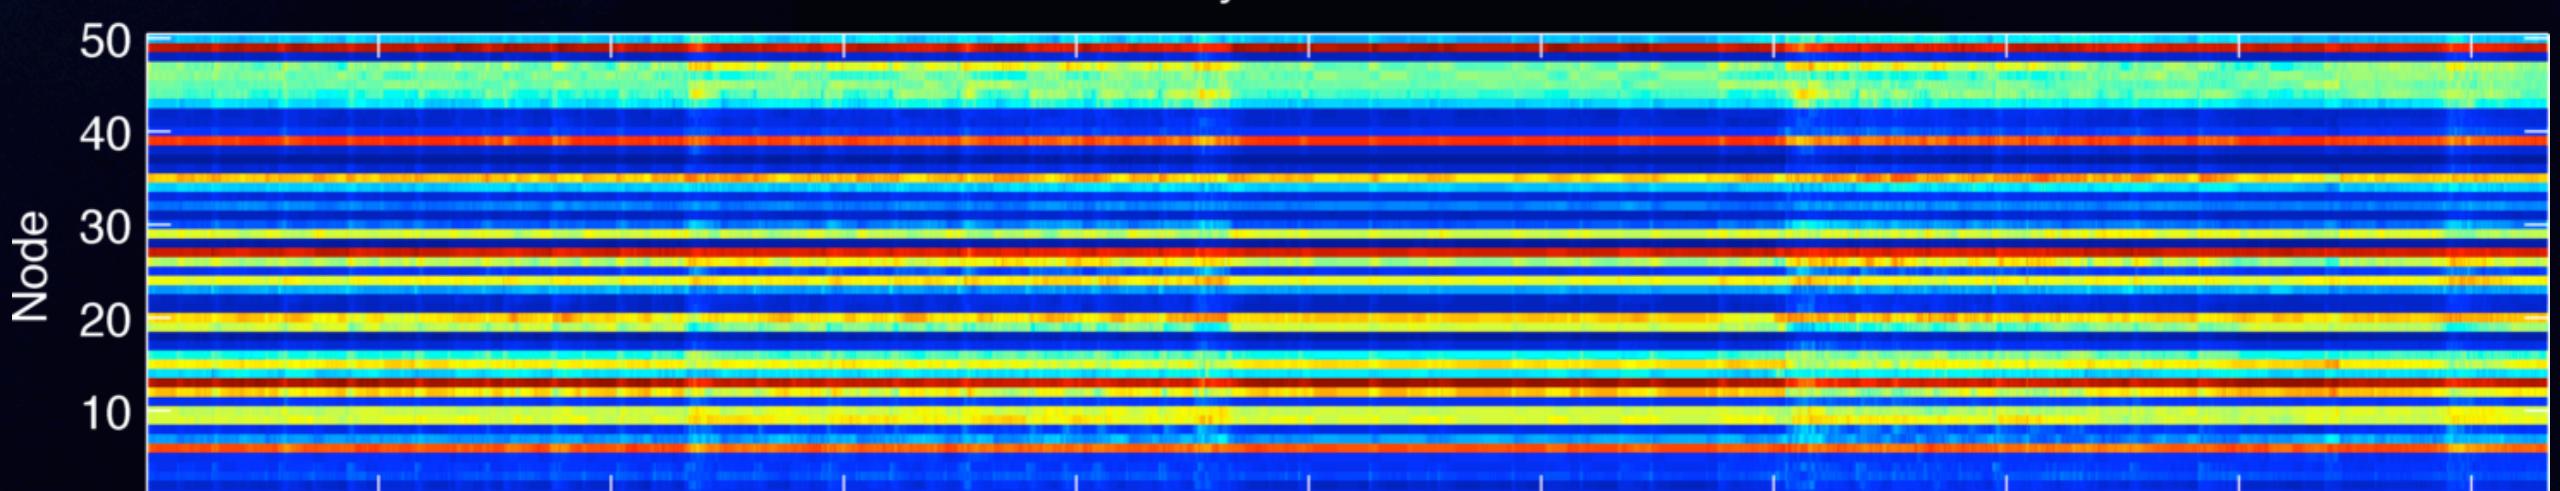
Magnitude Spectrogram  
Soulive: Come Together, 50-70 secs



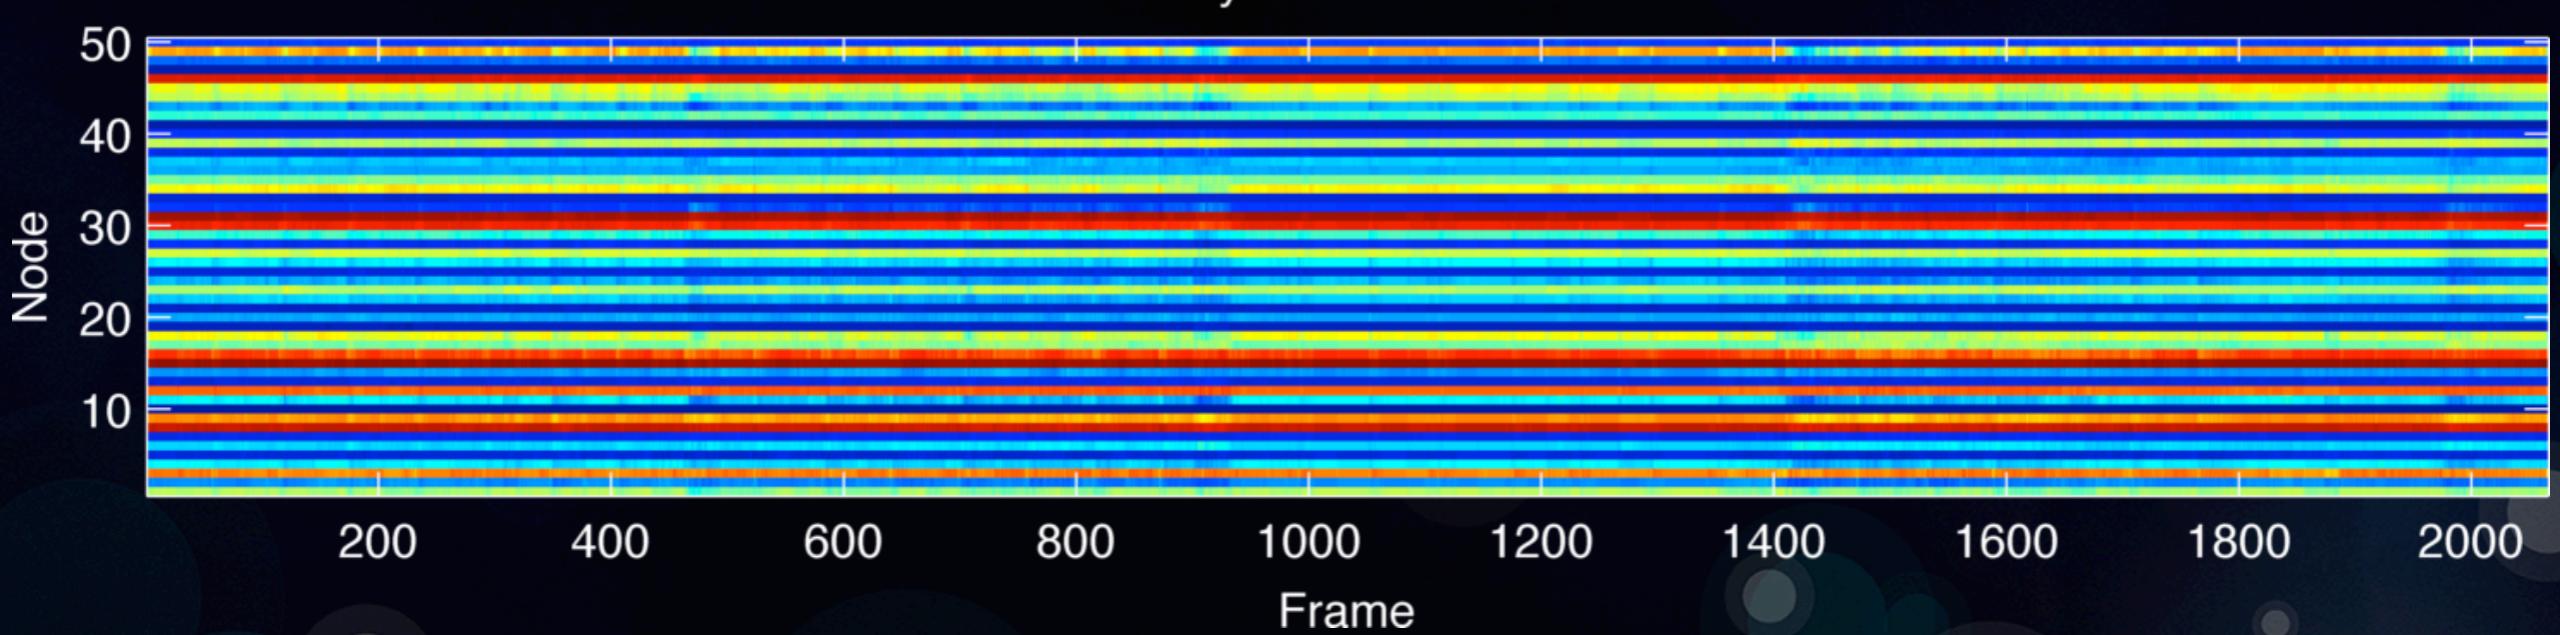
Layer 1 Features



Layer 2 Features



Layer 3 Features

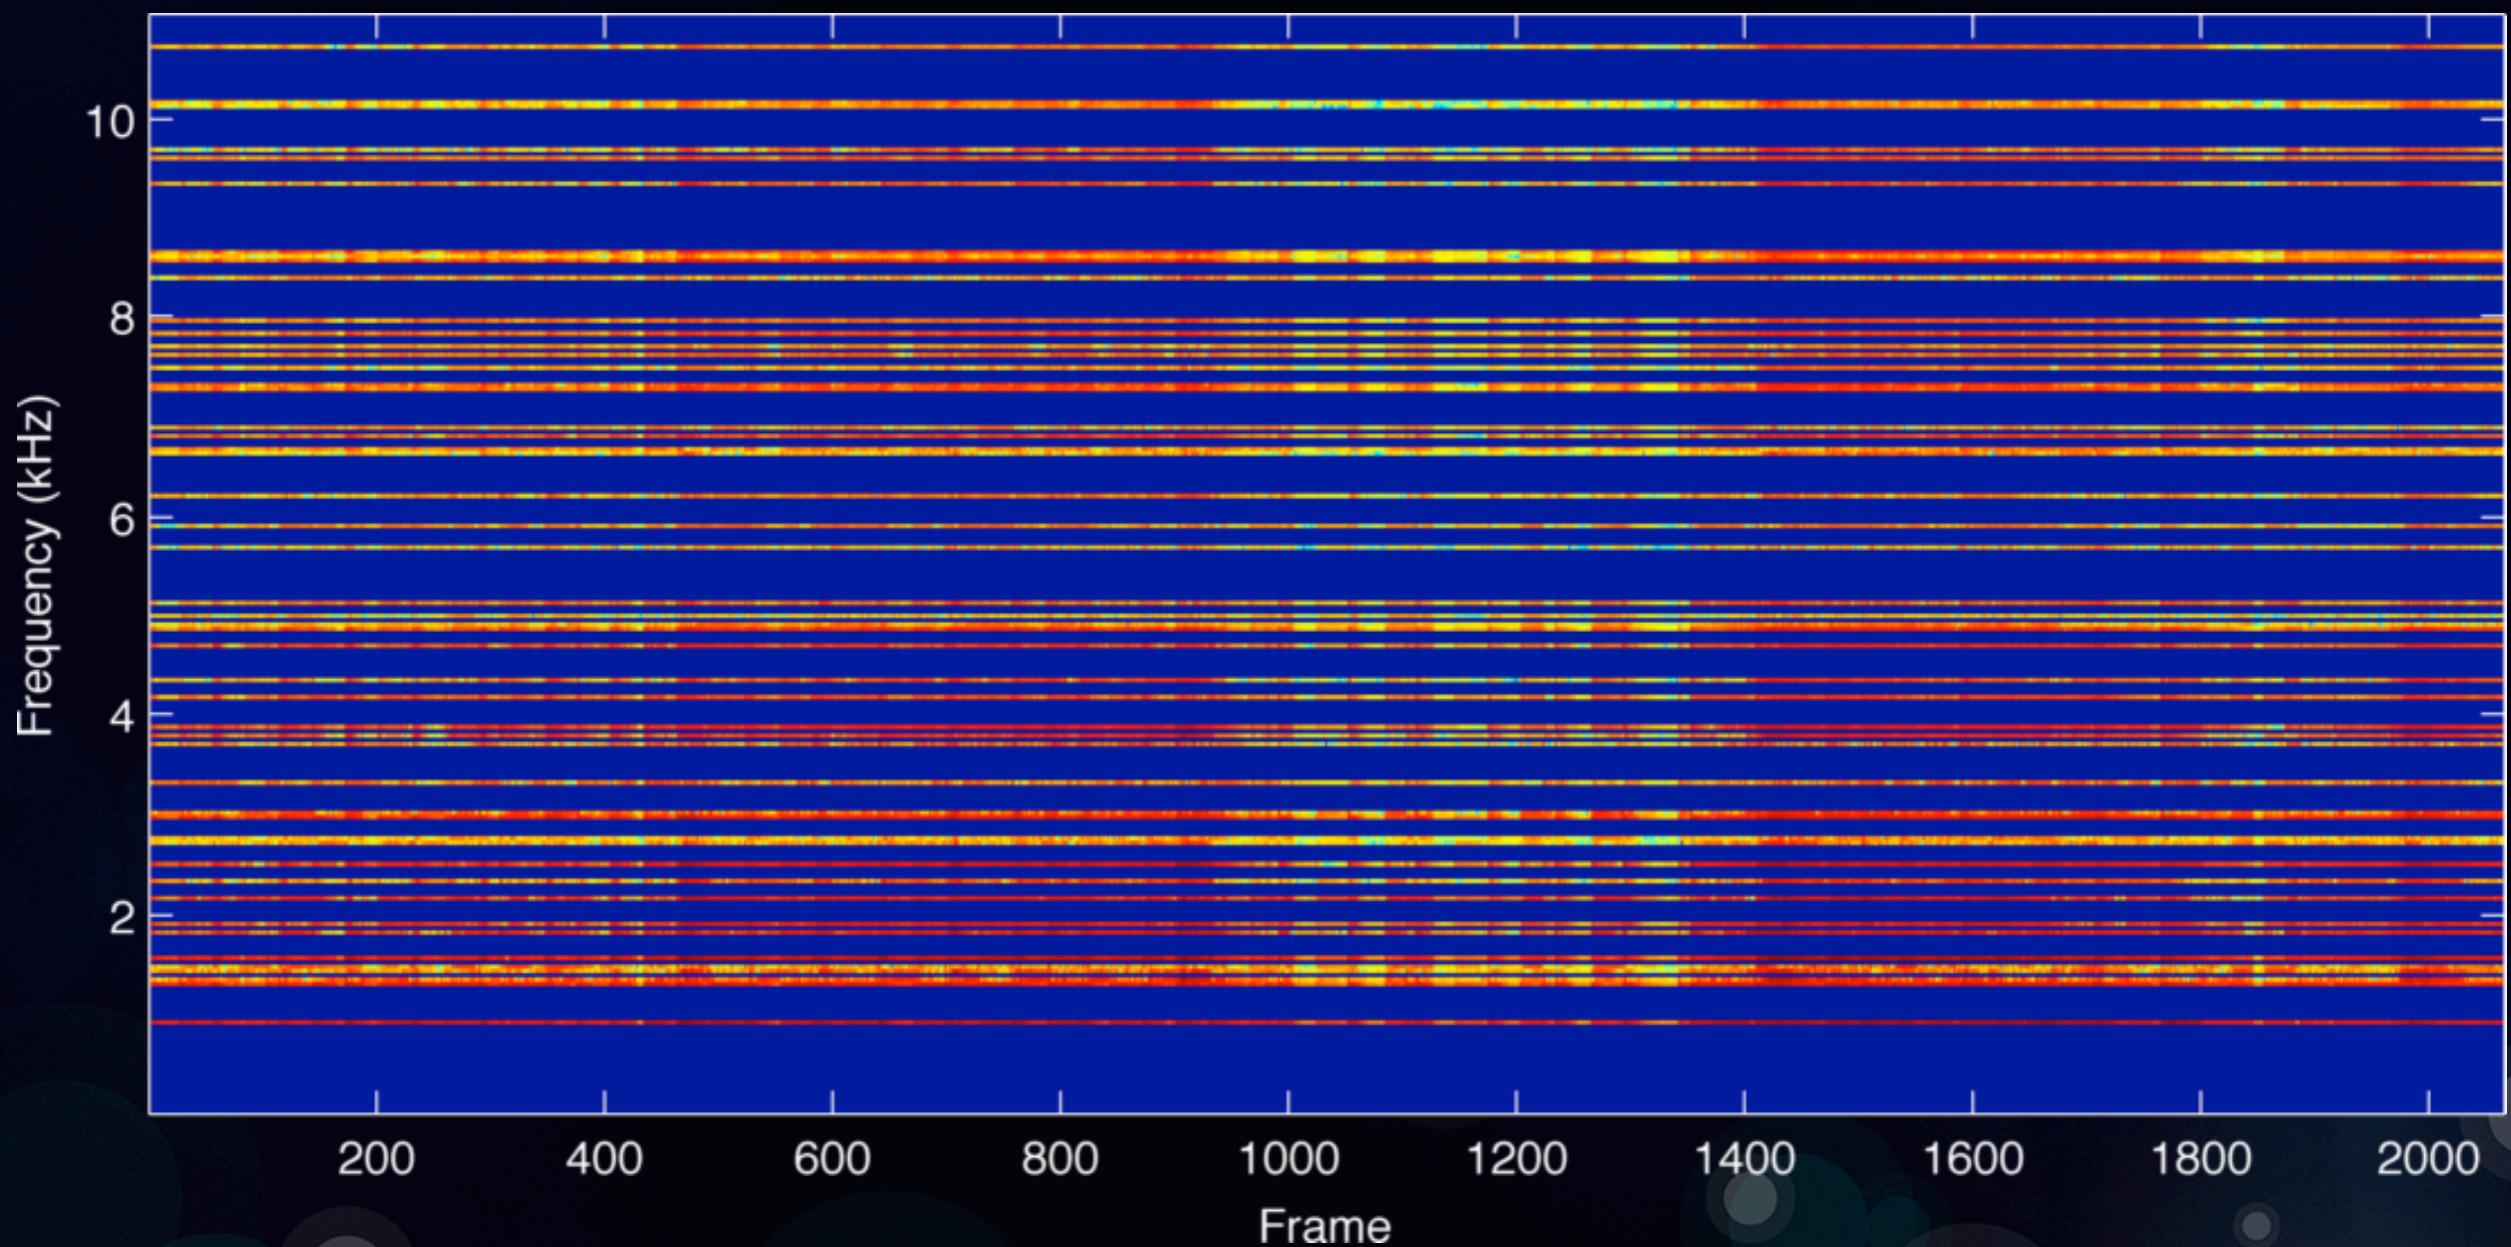


# STFT Reconstruction

$$\mathbf{y}_{\text{layer1}i} = \frac{1}{(1 + \exp(-\mathbf{x}_i \mathbf{w} - \mathbf{b}))}$$

# STFT Reconstruction

$$\mathbf{x}_{\text{reconstructed}i} = \left( -\ln \left( \frac{1}{\mathbf{y}_{\text{layer1}i}} - 1 \right) - b \right) \mathbf{w}^+$$



# Results

## Emotion prediction with learned features

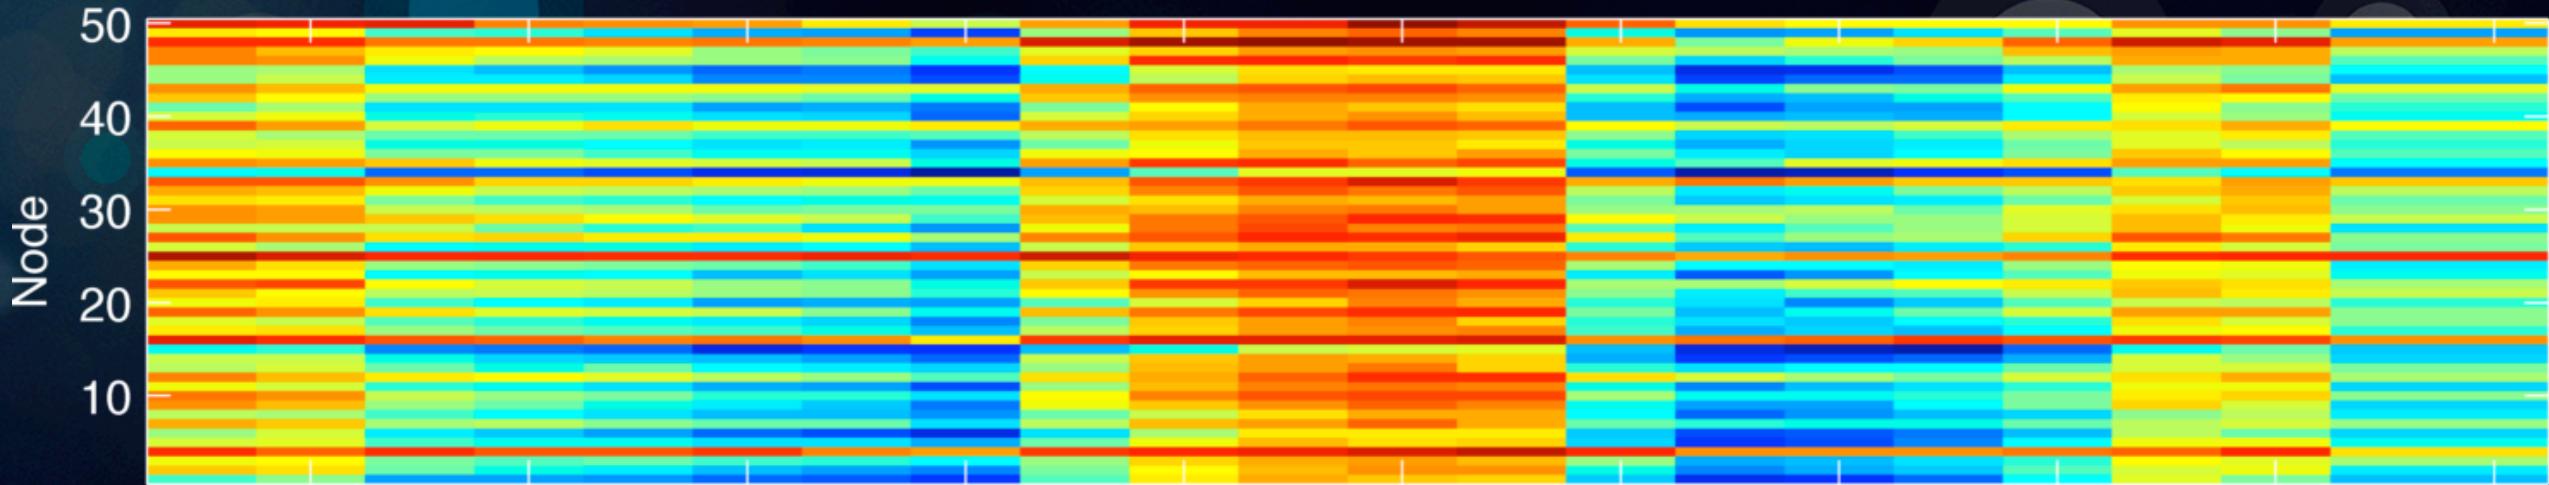
Feature/ Topology	Average Mean Distance	Average KL Divergence
MFCCs	$0.140 \pm 0.005$	$1.28 \pm 0.157$
Chroma	$0.182 \pm 0.006$	$3.33 \pm 0.294$
Spectral Shape	$0.153 \pm 0.006$	$1.51 \pm 0.160$
Spectral Contrast	$0.138 \pm 0.005$	$1.29 \pm 0.160$
ENTS	$0.151 \pm 0.006$	$1.41 \pm 0.175$
DBN Model Error	$0.203 \pm 0.009$	-
DBN Layer 1	$0.138 \pm 0.005$	$1.25 \pm 0.142$
DBN Layer 2	$0.133 \pm 0.004$	$1.19 \pm 0.129$
DBN Layer 3	$0.133 \pm 0.002$	$1.21 \pm 0.180$

# Next Steps

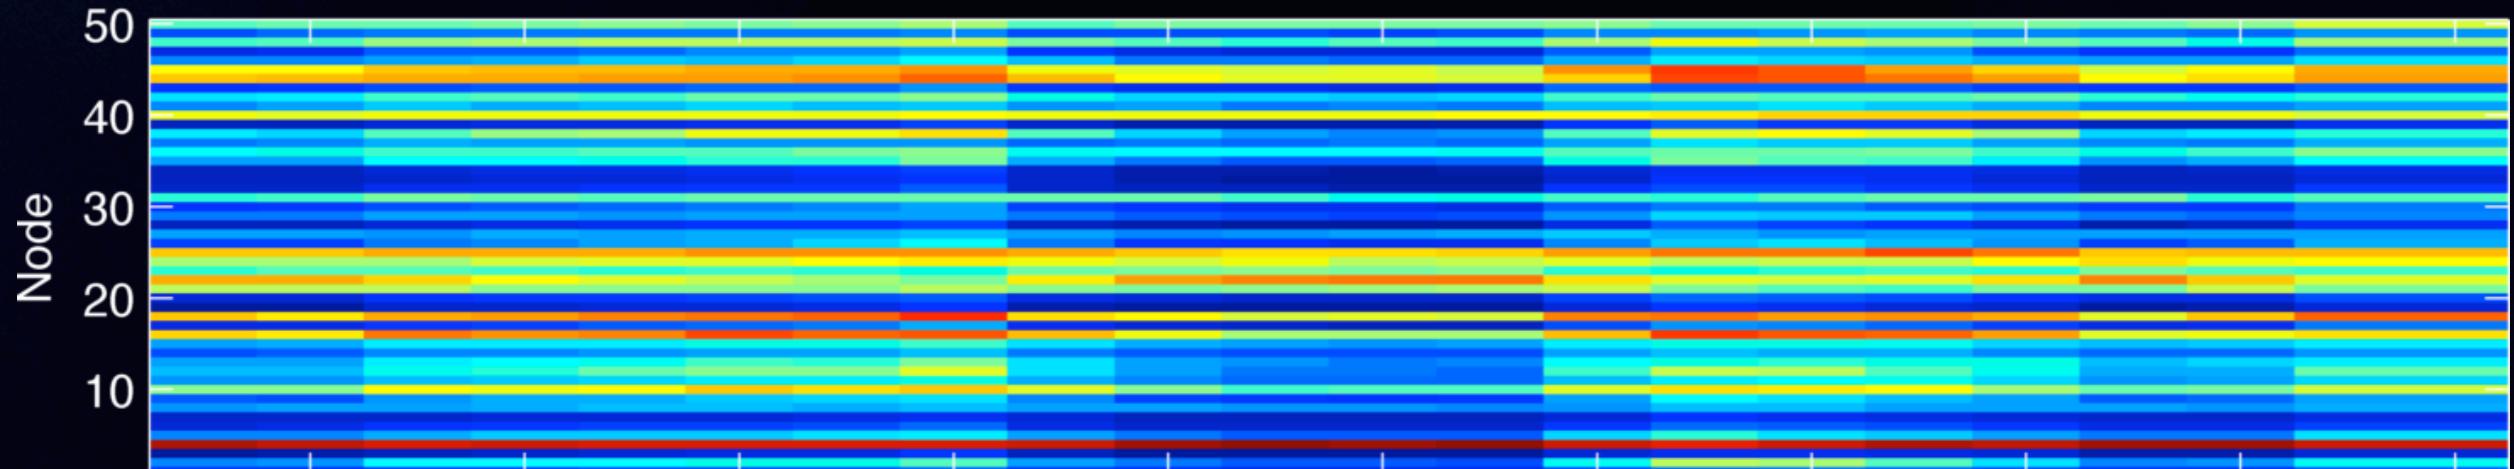
Developing more informed representations

- Current features are extracted from ~20 msec frames
  - No human can identify emotion at this time scale
  - Emotion is not expected to vary at this high a rate
- Instead look at multiple aggregated time scales
  - We investigate 1 sec, 2 sec, 4 sec scales combined

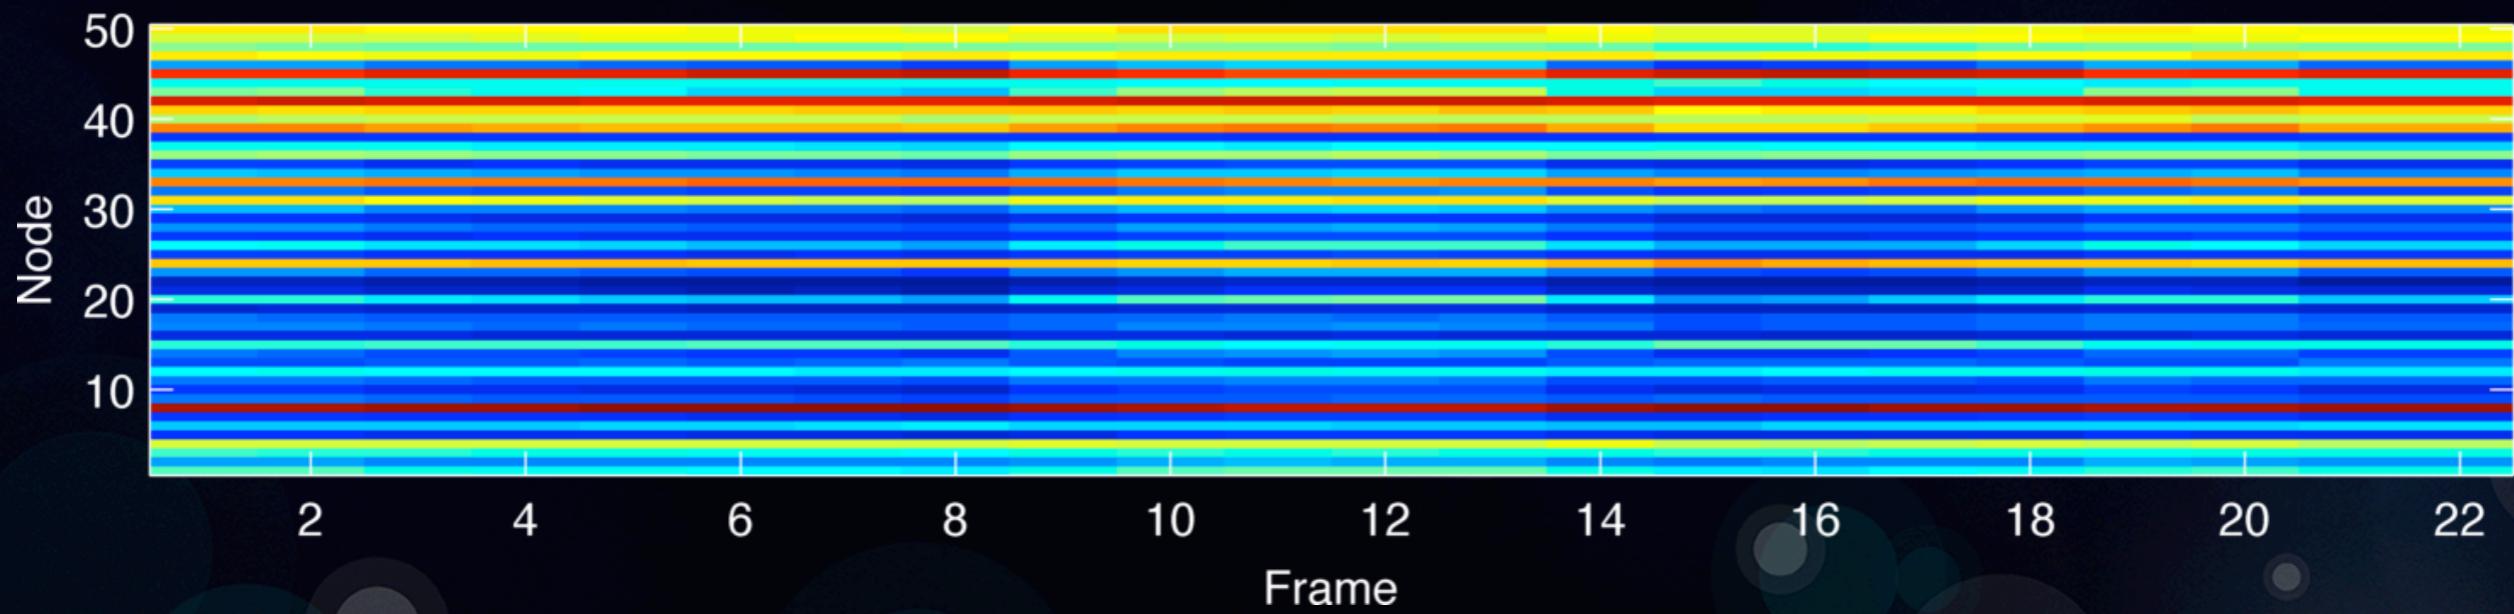
### Layer 1 Features



### Layer 2 Features



### Layer 3 Features



# Results

## Developing more informed representations

Feature/ Topology	Average Mean Distance	Average KL Divergence
DBN-SF Model Error	$0.203 \pm 0.009$	-
DBN-SF Layer 1	$0.138 \pm 0.005$	$1.25 \pm 0.142$
DBN-SF Layer 2	$0.133 \pm 0.004$	$1.19 \pm 0.129$
DBN-SF Layer 3	$0.133 \pm 0.002$	$1.21 \pm 0.180$
DBN-MF Model Error	$0.194 \pm 0.032$	-
DBN-MF Layer 1	$0.131 \pm 0.006$	$1.15 \pm 0.106$
DBN-MF Layer 2	$0.131 \pm 0.004$	$1.14 \pm 0.107$
DBN-MF Layer 3	$0.129 \pm 0.004$	$1.12 \pm 0.114$

# Harnessing Unlabeled Data

## Universal background model

- Boost performance by including much more data in the unsupervised phase
- No reason we should be limited to 240 song dataset
  - 50% training = 120 songs
  - Access to an additional ~7000 songs
  - ~1.5M examples
  - Develop generalized music model, fine-tune to predict emotion

# Results

## Universal background model

Feature/ Topology	Average Mean Distance	Average KL Divergence
DBN-MF Model Error	$0.194 \pm 0.032$	-
DBN-MF Layer 1	$0.131 \pm 0.006$	$1.15 \pm 0.106$
DBN-MF Layer 2	$0.131 \pm 0.004$	$1.14 \pm 0.107$
DBN-MF Layer 3	$0.129 \pm 0.004$	$1.12 \pm 0.114$
DBN-UBM Model Error	$0.140 \pm 0.015$	-
DBN-UBM Layer 1	$0.129 \pm 0.006$	$1.12 \pm 0.091$
DBN-UBM Layer 2	$0.128 \pm 0.004$	$1.13 \pm 0.097$
DBN-UBM Layer 3	$0.128 \pm 0.004$	$1.11 \pm 0.090$

# Selecting Input Representations

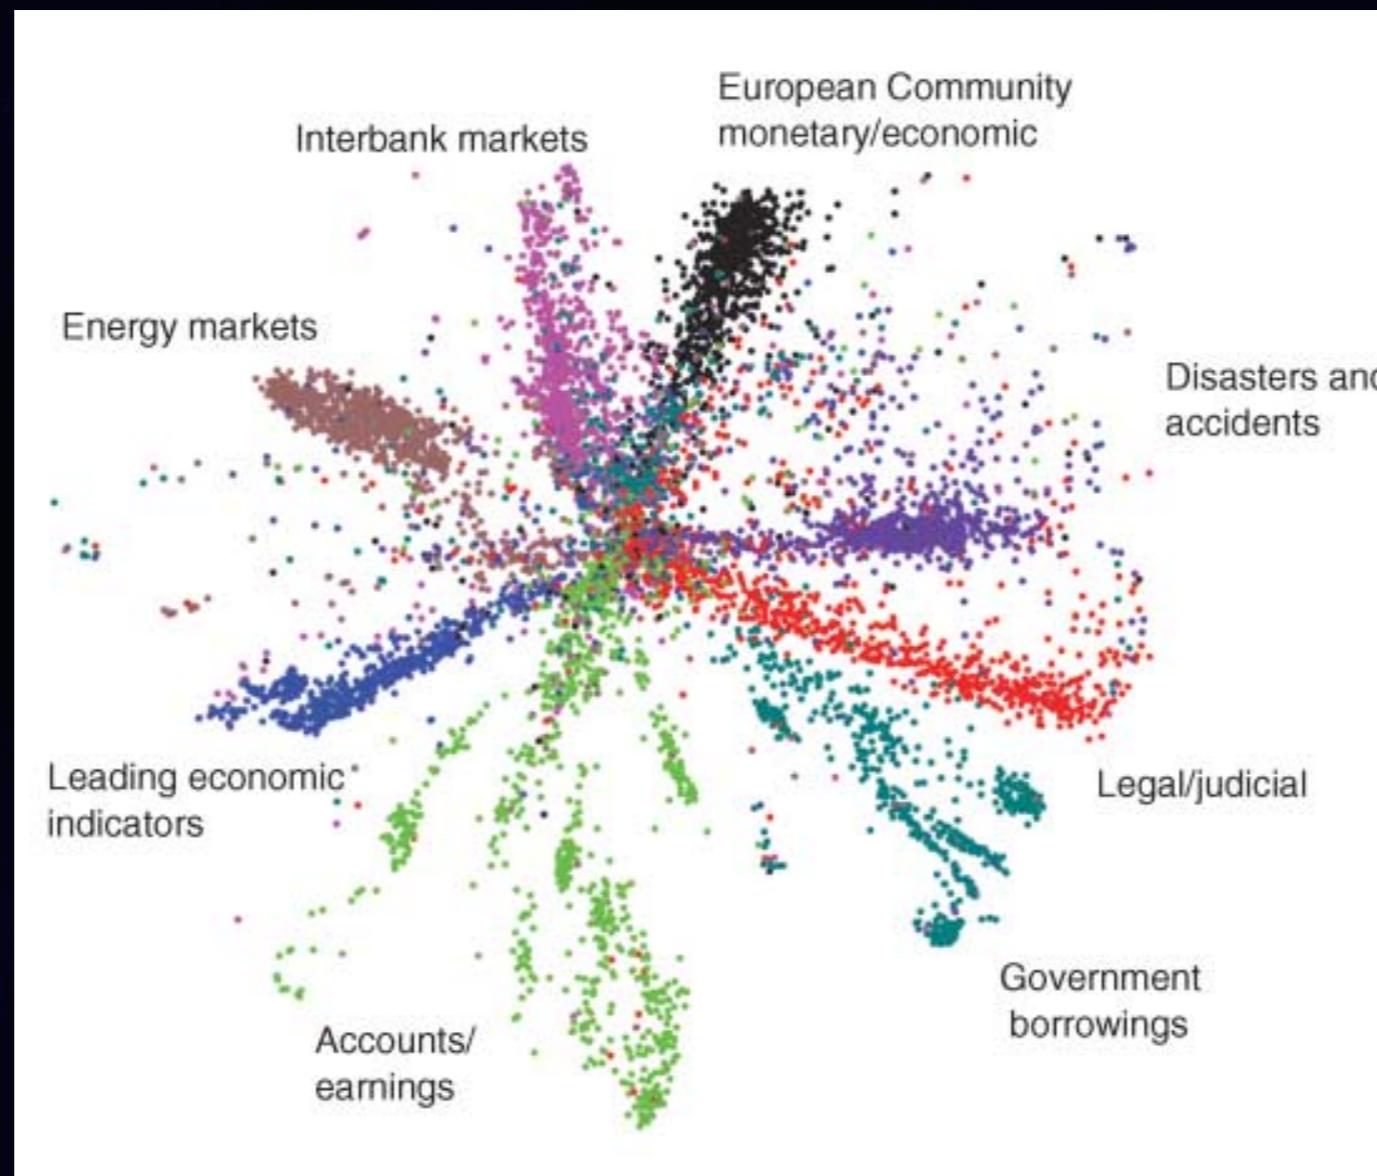
Taking best advantage of pre-training

- Work thus far used **magnitude spectra** as the input representation
  - Sought direct mapping between acoustic data and emotion space parameters
  - Saw **modest** performance **gains**
- Unsupervised methods tend to highlight the most prominent aspects of data
  - Such clustering with spectrograms may not provide a useful starting point for gradient descent

# Learning Deep Autoencoders

## Stacking RBMs

2000-500-250-125-2 Autoencoder (Figure 4, C)



G. E. Hinton and R. R. Salakhutdinov (2006). Reducing the Dimensionality of Data with Neural Networks. *Science*.

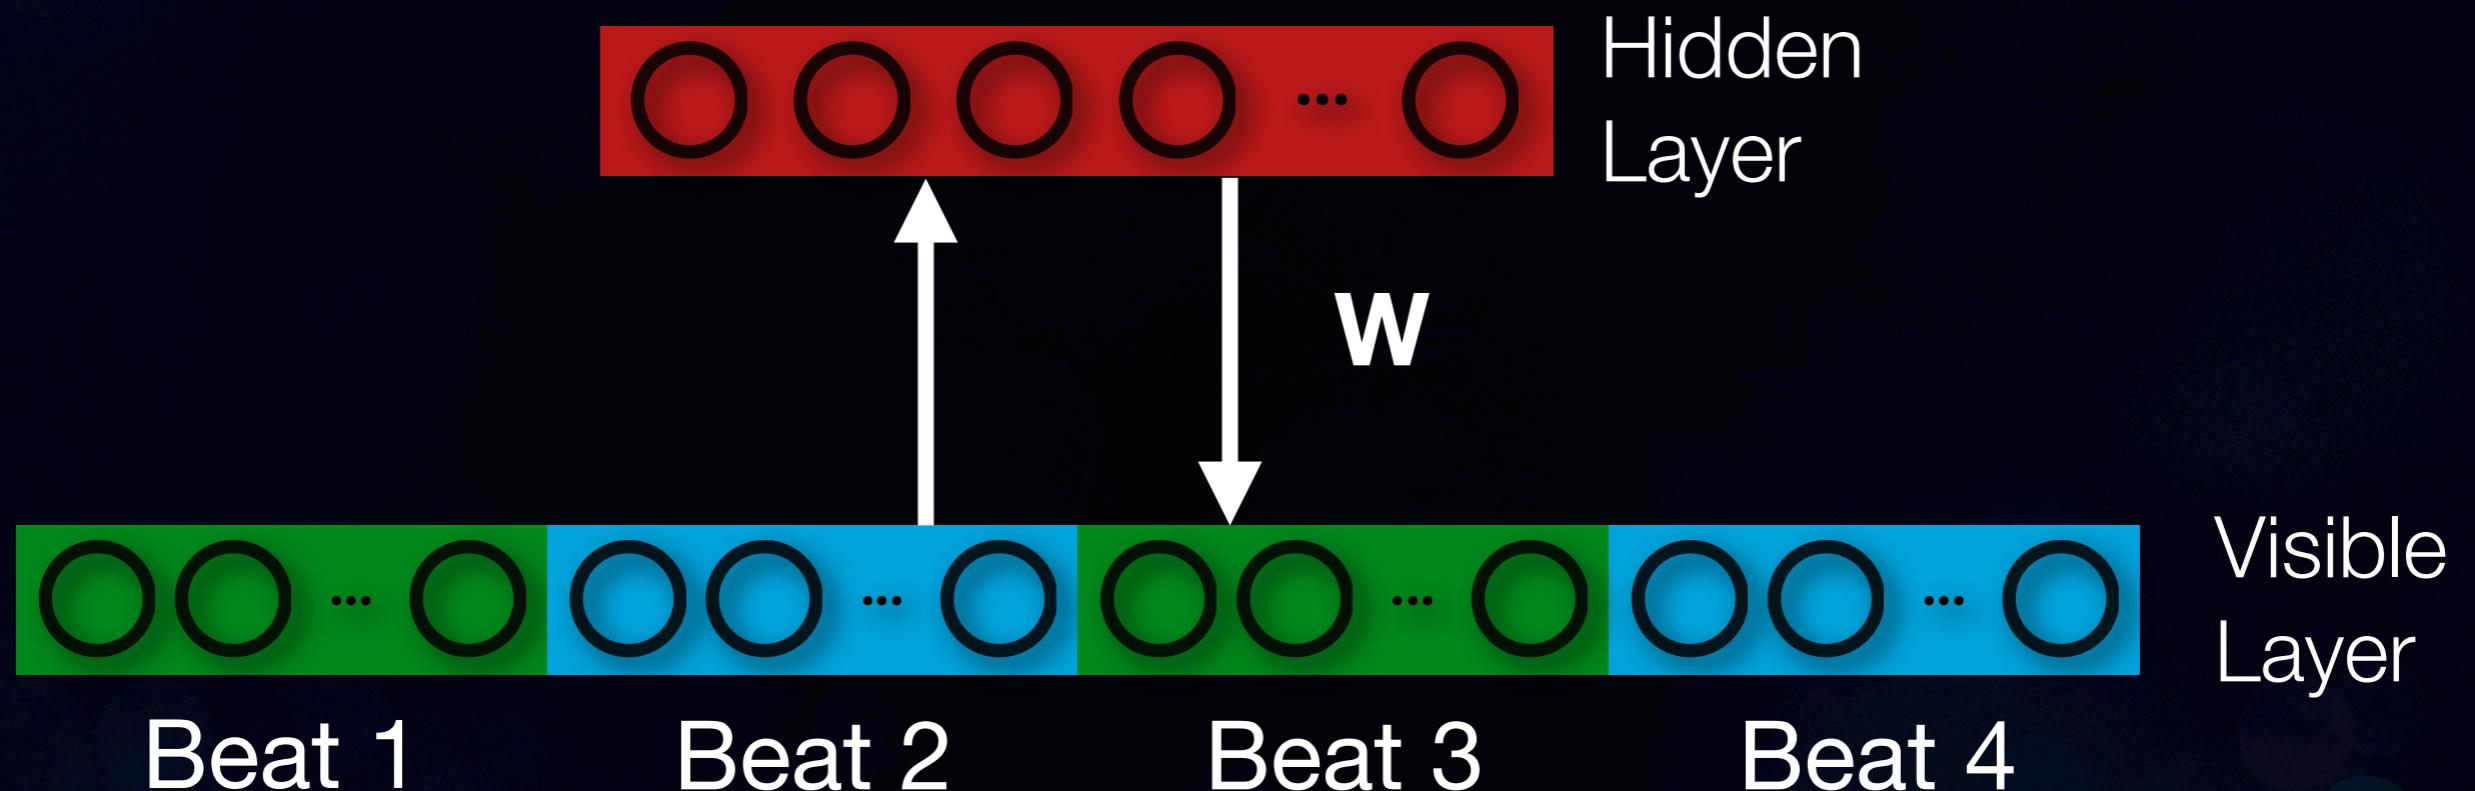
# Rhythm and Melody DBNs

Learning rhythmic patterns and melodic contour

- Rhythm and melody are musical quantities that are understood to inform emotion
- Learn informative domains from pre-training alone
  - Harmonic percussive/source separation
  - Parse rhythmic from melodic contributions
  - Beat synchronous representations
  - Parse musical events

# Going Even Further

Learning rhythmic patterns and melodic contour



# MediaEval Dataset

1000 songs for music emotion recognition

- 1000 songs for music emotion recognition
  - All songs selected from the Free Music Archive
    - Creative commons licensed
  - Data is annotated using Amazon's Mechanical Turk
    - For each song a 40-second clip is analyzed
    - Each song is analyzed by a minimum of 10 participants
    - Continuous and discrete Valence-Arousal is collected

# Thank You!

[eschmidt@pandora.com](mailto:eschmidt@pandora.com)